

Bushwhacking your way around a bootloader

Rebecca ".bx" Shapiro
2018.11.16



Tools and techniques for traversing treacherous code bases
-or- How I managed to develop understanding of U-Boot
Blackhoodie Berlin

whoami

Dr. .bx

- Senior security researcher @ Narf Industries
- Studied w/ Sergey Bratus & the Dartmouth Trust Lab
- Commander of ELF metadata-based weird machines
- ELF, bootloaders,
- Dynamic analysis
- Defensive research (more or less) with a dash of reverse engineering
- [at]bxsays on Twitter

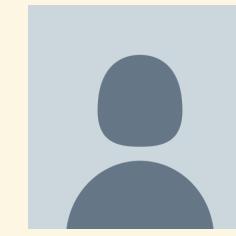


Photo circa Sept 2018



Meet Das U-Boot bootloader

```
U-Boot 2010.12-xes_r3 (Aug 25 2011 - 11:04:04)

CPU0: P2020E, Version: 1.0, (0x80ea0010)
Core: E500, Version: 4.0, (0x80211040)
Clock Configuration:
    CPU0:1066.667 MHz, CPU1:1066.667 MHz,
    CCB:533.333 MHz,
    DDR:400 MHz (800 MT/s data rate) (Asynchronous), LBC:133.333 MHz
L1: D-cache 32 kB enabled
    I-cache 32 kB enabled
Board: X-ES XPedite5501 PMC/XMC SBC
    Rev SA, Serial# 36093001, Cfg 90015130-1
I2C: ready
DRAM: 2 GiB (DDR3, 64-bit, CL=6, ECC on)
FLASH: Executed from FLASH1
POST memory PASSED
FLASH: 256 MiB
L2: 512 KB enabled
NAND: 4096 MiB
PCIE1: connected as Root Complex (no link)
PCIE1: Bus 00 - 00
PCIE2: disabled
PCIE3: disabled
In: serial
Out: serial
Err: serial
DTT: 37 C local / 59 C remote (adt7461@4c)
DTT: 37 C local (lm75@48)
Net: eTSEC1 connected to Broadcom BCM5482S
eTSEC2 connected to Broadcom BCM5482S
eTSEC1, eTSEC2
POST i2c PASSED
Hit any key to stop autoboot: 0
=> |
```

Meet Das U-Boot bootloader

Language	files	blank	comment	code
C	3958	177722	230606	911861
C/C++ Header	3540	64684	108111	429854
Assembly	236	5927	10632	24037
Python	119	4380	9180	12486
Perl	6	1660	1346	9850
make	911	2263	4664	8500
Bourne Shell	32	427	626	2164
C++	1	233	58	1588
yacc	2	169	75	1076
Glade	1	58	0	603
lex	2	98	41	539
NAnt script	1	91	0	367
YAML	1	13	25	347
Bourne Again Shell	3	75	66	316
Markdown	1	80	0	283
DOS Batch	3	20	0	176
CSS	2	24	10	90
Kermit	3	4	20	83
Tcl/Tk	1	5	5	28
sed	2	1	27	24
INI	2	3	0	14
XSLT	1	0	1	9
SUM:	8828	257937	365493	1404295

Meet Das U-Boot bootloader

```
[user@boot-dev ~]$ cloc u-boot/
 13518 text files.
 12700 unique files.
 4701 files ignored.
github.com/AlDanial/cloc v 1.76 T=4.02 s (2196.7 files/s, 504571.1 lines/s)
```

Language	files	blank	comment	code
C	3958	177722	230606	911861
C/C++ Header	3540	64684	108111	429854
Assembly	236	5927	10632	24037
Python	119	4380	9180	12486
Perl	6	1660	1246	9850
makefile	1	1660	1346	9850
Bourne Shell	32	427	626	2164
C-Header	1	288	58	1588
YAML	2	58	75	1076
Glade	1	0	0	603
JPEG	2	168	41	539
Nano	1	169	0	367
YAML	1	58	25	347
Bourne Again Shell	1	75	66	316
Makefile	1	80	0	283
DOS Batch	3	20	0	176

"Only" 111 MB of code for a resource-constrained system's bootloader

```
[user@boot-dev ~]$ make -C u-boot distclean
make: Entering directory '/home/user/u-boot'
make: Leaving directory '/home/user/u-boot'
[user@boot-dev ~]$ rm -rf u-boot/.git
[user@boot-dev ~]$ du -sh u-boot/
111M u-boot/
```

CSS	2	24	10	90
Kermit	3	4	20	83
Tcl/Tk	1	5	5	28
sed	2	1	27	24
INI	2	3	0	14
XSLT	1	0	1	9

SUM:	8828	257937	365493	1404295
------	------	--------	--------	---------

Meet Das U-Boot bootloader

```
[user@boot-dev ~]$ cloc u-boot/
 13518 text files.
 12700 unique files.
 4701 files ignored.
github.com/AlDanial/cloc v 1.76 T=4.02 s (2196.7 files/s, 504571.1 lines/s)
```

Language	files	blank	comment	code
C	3958	177722	230606	911861
C/C++ Header	3540	64684	108111	429854
Assembly	236	5927	10632	24037
Python	119	4380	9180	12486
Perl	6	1660	1246	9850

"Only" 111 MB of code for a resource-constrained system's bootloader

```
[user@boot-dev ~]$ make -C u-boot distclean
make: Entering directory '/home/user/u-boot'
make: Leaving directory '/home/user/u-boot'
[user@boot-dev ~]$ rm -rf u-boot/.git
[user@boot-dev ~]$ du -sh u-boot/
111M  u-boot/
```

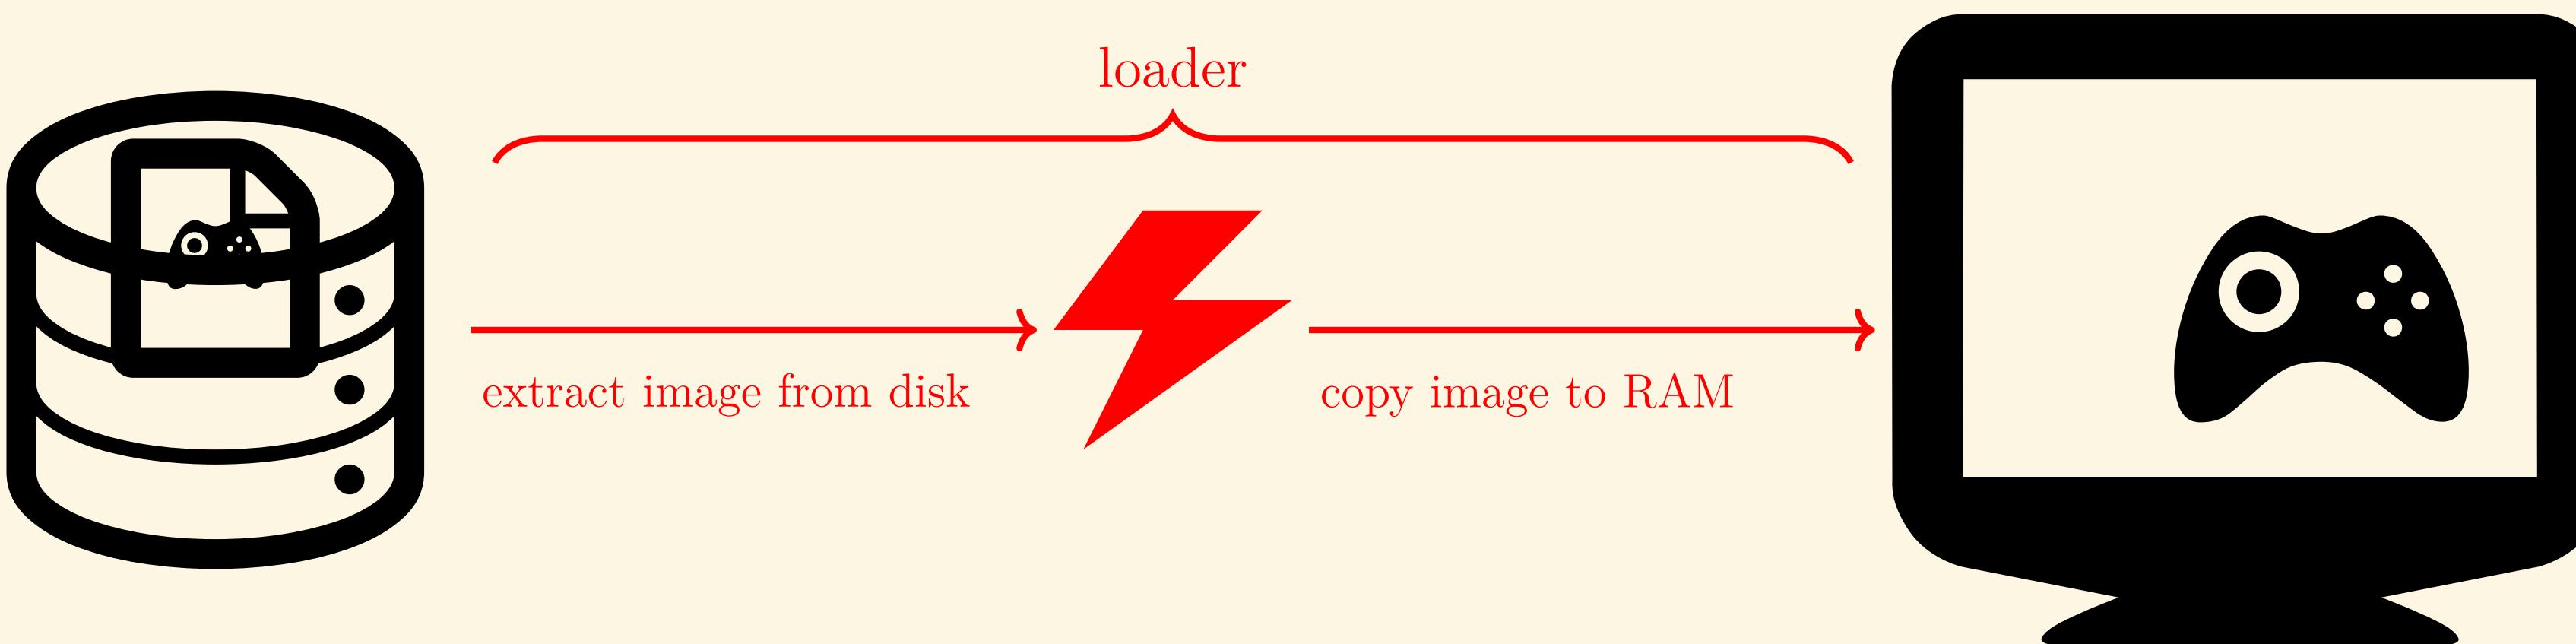
CSS	2	24	10	90
Kermit	3	4	20	83
Tcl/Tk	1	5	5	28
sed	2	1	27	24
INI	2	3	0	14
XSLT	1	0	1	9

SUM:	8828	257937	365493	1404295
------	------	--------	--------	---------

Quick aside: what is a loader?

The magic that transforms a binary image into an running application

- **Loader:** Software that transduces binary images into memory for execution
- **Binary image:** Static representation/encapsulation of binary (machine) code
 - e.g. An ELF or PE file



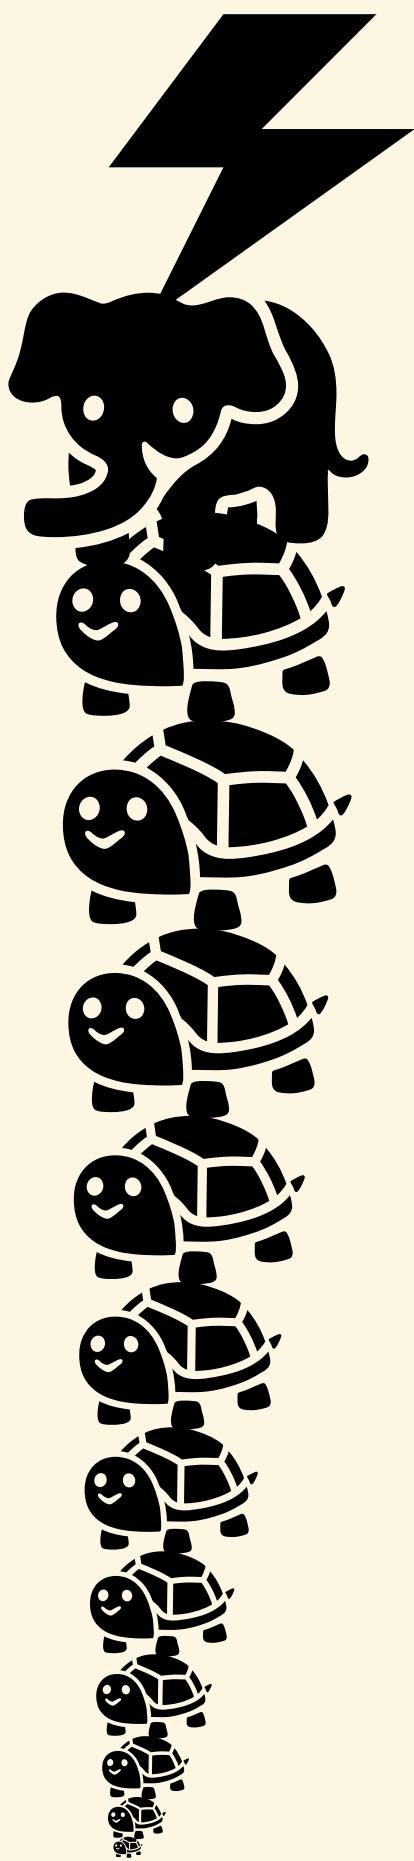
Other useful terminology

- **Address space:** general term referring to addressable memory
- **Memory map:** address space model that semantically labels memory regions

Who loads the loader?

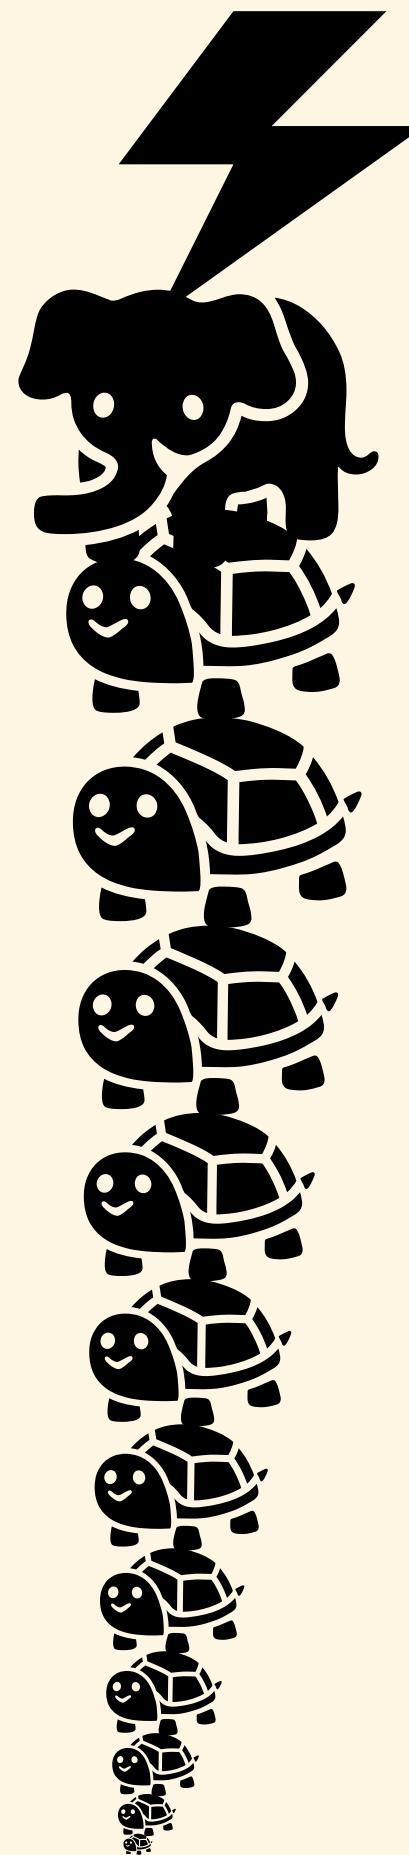
A loader, of course

(It's turtles all the way down)

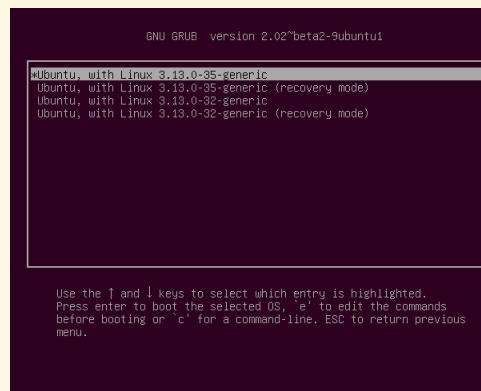


Who loads the loader?

A loader, of course
(It's turtles all the way down)

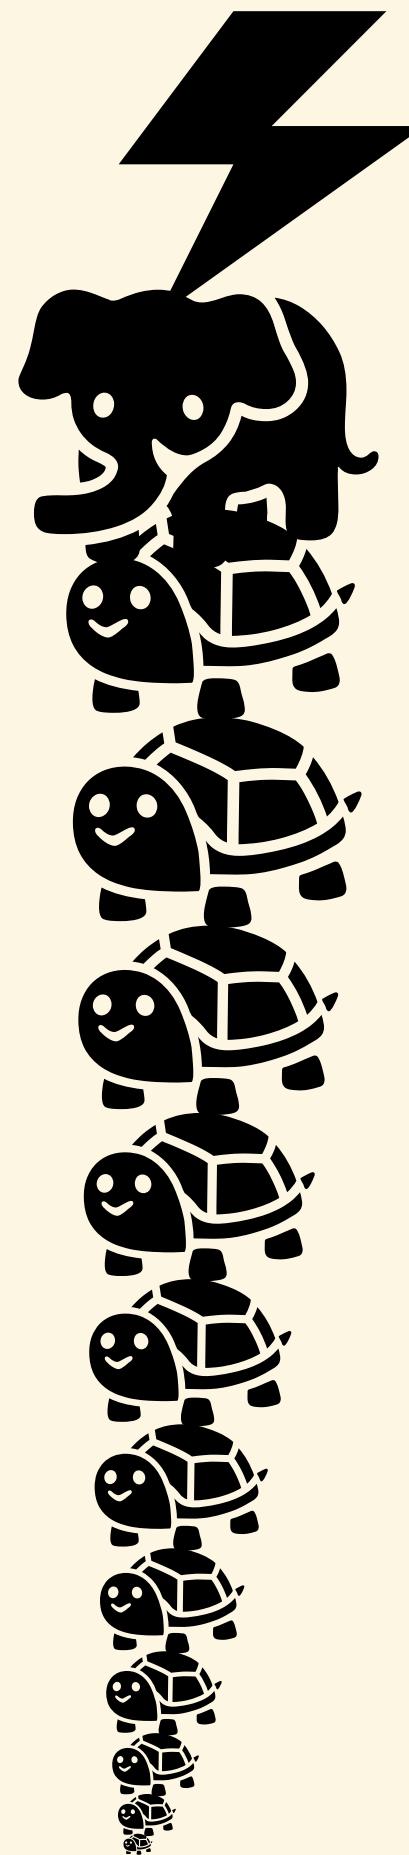


Bootloaders: a subset of loaders
that execute before the OS (or primary
application) is executed

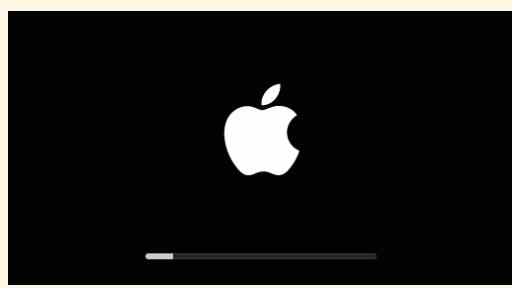


Who loads the loader?

A loader, of course
(It's turtles all the way down)



Bootloaders: a subset of loaders
that execute before the OS (or primary
application) is executed



</aside>

The existential question.



Overall research goals

1. Identify **weaknesses** underlying (boot)loader security
2. Develop **(boot)loader hardening techniques** that:
 - are **realistic**
 - lend themselves to **formal reasoning**
 - can be **retroactively applied to existing** loaders
3. Demonstrate technique **feasibility**

The existential question.



Overall research goals

1. Identify **weaknesses** underlying (boot)loader security
2. Develop **(boot)loader hardening techniques** that:
 - are **realistic**
 - lend themselves to **formal reasoning**
 - can be **retroactively applied** to **existing** loaders
3. Demonstrate technique **feasibility**

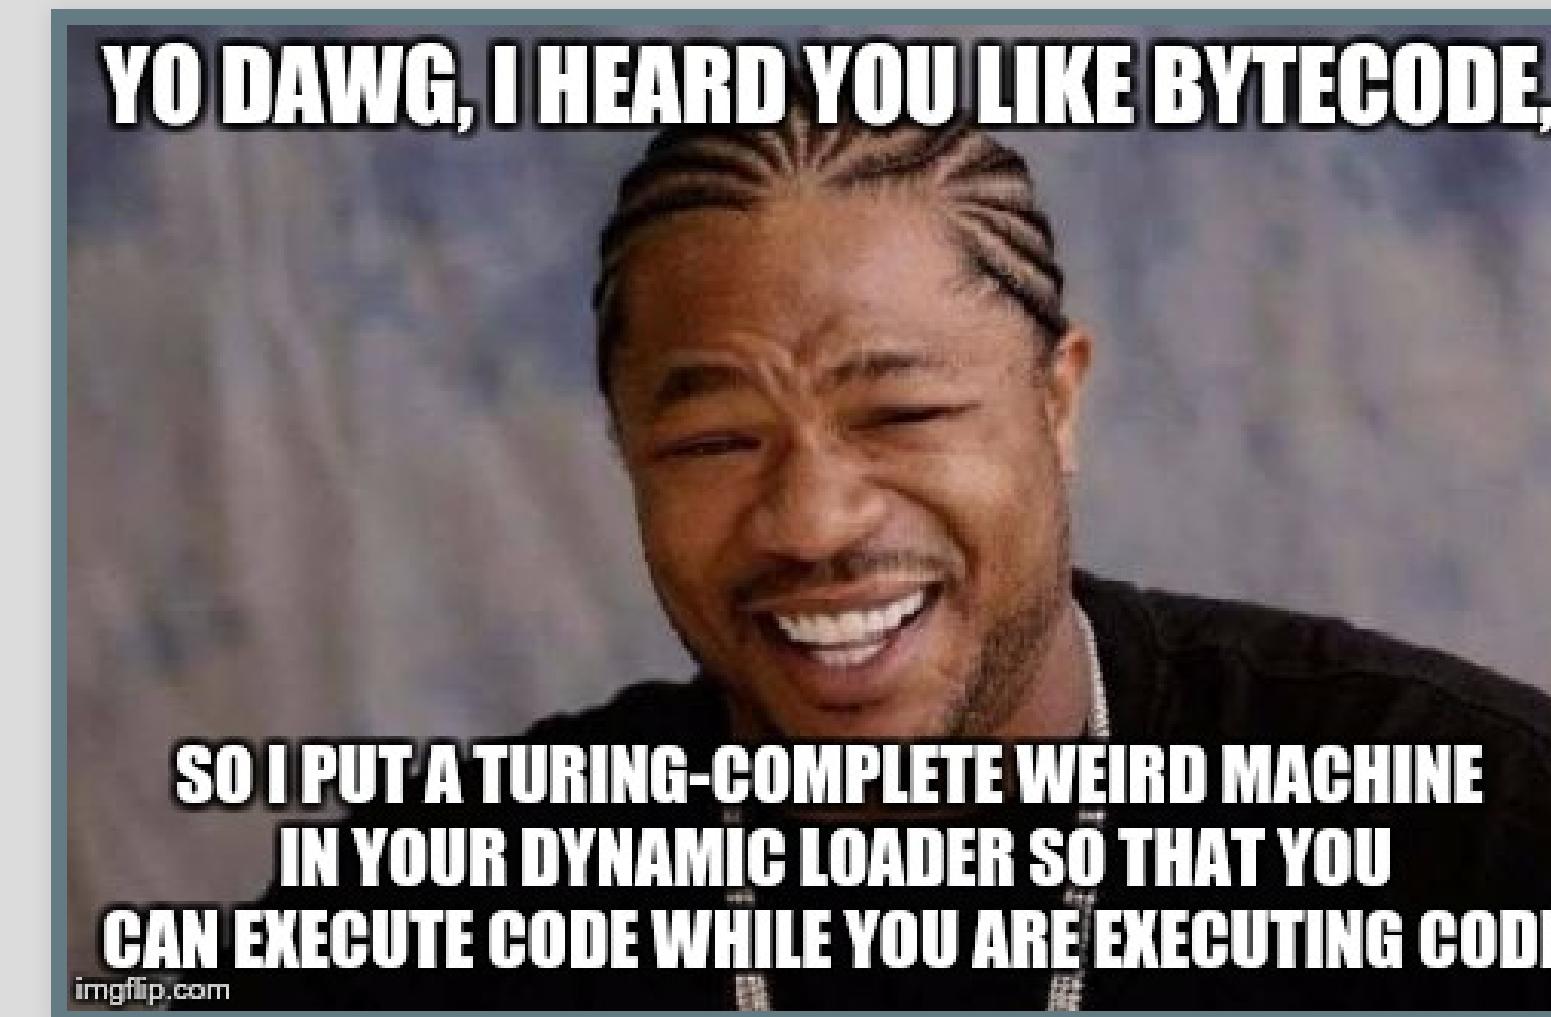
Hence the bootloader reversing

The existential question.

* *Flashback to 2012* *

The Turing-complete ELF-metadata weird machine

a brainfuck to ELF-metadata compiler @ <https://github.com/bx/elf-bf-tools>



(ELF is *NIX's file format for executables, libraries, etc.)



DEF CON, 29C3, USENIX WOOT, PoC || GTFO

[\[1\]](#), [\[2\]](#), [\[3\]](#), [\[4\]](#)

Hence the bootloader reversing

The existential question.



Overall research goals

1. Identify **weaknesses** underlying (boot)loader security
2. Develop **(boot)loader hardening techniques** that:
 - are **realistic**
 - lend themselves to **formal reasoning**
 - can be **retroactively applied** to **existing** loaders
3. Demonstrate technique **feasibility**

Hence the bootloader reversing

The existential question.

BUT WHY?

* The *ultimate* goal *



1. Identify
2. Develop
 - are real
 - lend them
 - can be re
3. Demonstra

security
es that:

loaders

Hence the bootloader reversing

The existential question.



Overall research goals

1. Identify **weaknesses** underlying (boot)loader security
2. Develop **(boot)loader hardening techniques** that:
 - are **realistic**
 - lend themselves to **formal reasoning**
 - can be **retroactively applied** to **existing** loaders
3. Demonstrate technique **feasibility**

Hence the bootloader reversing

This talk for those in a hurry



- ~~Introduce goals and case study~~
- Pontificate about bootloaders in general
- Debugging U-Boot (as according to U-Boot)
- My instrumentation toolsuite
 - An attempt at something better
- Techniques for identifying code <-> data
- A test drive of a simple example

Properties of a bootloader

- Load images and prepare its address space
- Initialize resources/hardware
- Sometimes self-relocate

In general, (boot)loaders:

- Allocate **non-overlapping** addresses
- Manage address **alignment requirements**
- **Prepare memory map** for target
 - **Load** target image into memory
 - **Patch** (link) loaded images
- **Extract** and **enforce** requirements and restrictions imposed by both resources and target



A bit about self-relocation

Bootloaders self-relocate, this may sound crazy
but when RAM is small, you can't be lazy

```
ENTRY(relocate_code)
    ldr r1, =__image_copy_start /* r1 <- &__image_copy_start */
*/
    subs    r4, r0, r1      /* r4 <- relocation offset */
    beq relocate_done        /* skip relocation */
    ldr r2, =__image_copy_end /* r2 <- &__image_copy_end */
copy_loop:
    ldmia  r1!, {r10-r11} /* copy from source address [r1]
*/
    stmia  r0!, {r10-r11} /* copy to target address [r0]
*/
    cmp r1, r2      /* until source end address [r2] */
    blo copy_loop
    /* fix .rel.dyn relocations */
    ldr r2, =__rel_dyn_start /* r2 <- SRC &__rel_dyn_start */
    ldr r3, =__rel_dyn_end /* r3 <- SRC &__rel_dyn_end */
fixloop:
    ldmia  r2!, {r0-r1} /* (r0,r1) <- (SRC location,fixup) */
    and r1, r1, #0xff
    cmp r1, #23 /* relative fixup? */
    bne fixnext
    /* relative fix: increase location by offset */
    add r0, r0, r4
    ldr r1, [r0]
    add r1, r1, r4
    str r1, [r0]
fixnext:
    cmp r2, r3
    blo fixloop
relocate_done:
    bx lr
ENDPROC(relocate_code)
```



Questions to answer when you are exploring a bootloader

- On what **architecture** does it run? (ARM? x86?)
- **Where** can it be **stored**? (First disk sector? Flash memory?)
- **Where** in memory can it be **loaded**?
- **How** is its image **packaged**? (MBR? ELF?)
- **What** runtime **arguments** does it expect and how does it **locate** them?
- From **what devices** can it load its **next stage**? (Disk? Network?)
- **What formats** can it handle as it locates the next stage? (FAT? GPT?)
- How can the **next stage** be **packaged**? (ELF? Compressed? Signed?)
- **What arguments** does it pass to the next stage and how?



Phase one

Let's begin exploring U-Boot



Bootloader instrumentation: tricky bits

1. Relocation
2. Magic/hardcoded numbers
3. Major characteristics of bootloader's address space may change
4. Possible "holes" in its address space
5. Operations or peripherals may silently fail (e.g., nothing may happen if a read-only register has been written)
6. Incorrect documentation-based bugs
7. Operating mode of CPU may change (e.g., privileged, unprivileged)



Bootloader instrumentation: tricky bits

1. Relocation
2. Magic/hardcoded numbers
3. Major characteristics of bootloader's address space may change
4. Possible "holes" in its address space
5. Operations or peripherals may silently fail (e.g., nothing may happen if a read-only register has been written)
6. Incorrect documentation-based bugs
7. Operating mode of CPU may change (e.g., privileged, unprivileged)



Bootloader instrumentation in action

(with GDB and QEMU)

```
(gdb) target remote | qemu-system-arm -gdb stdio -M beaglebm -sd sd.img -S
```

Bootloader instrumentation in action

(with GDB and QEMU)

```
(gdb) target remote | qemu-system-arm -gdb stdio -M beaglebm -sd sd.img -S  
Remote debugging using | qemu-system-arm -gdb stdio -M beaglebm -sd sd.img -S  
0x40014000 in ?? ()
```

Bootloader instrumentation in action

(with GDB and QEMU)

```
(gdb) target remote | qemu-system-arm -gdb stdio -M beaglebm -sd sd.img -S  
Remote debugging using | qemu-system-arm -gdb stdio -M beaglebm -sd sd.img -S  
0x40014000 in ?? ()  
(gdb) continue  
Continuing.
```

Bootloader instrumentation in action

(with GDB and QEMU)

```
(gdb) target remote | qemu-system-arm -gdb stdio -M beaglexm -sd sd.img -S  
Remote debugging using | qemu-system-arm -gdb stdio -M beaglexm -sd sd.img -S  
0x40014000 in ?? ()  
(gdb) continue  
Continuing.  
Breakpoint 1, jump_to_image_no_args (spl_image=0x80000000 )  
at arch/arm/cpu/armv7/omap-common/boot-common.c:229
```

Bootloader instrumentation in action

(with GDB and QEMU)

```
(gdb) target remote | qemu-system-arm -gdb stdio -M beaglexm -sd sd.img -S
Remote debugging using | qemu-system-arm -gdb stdio -M beaglexm -sd sd.img -S
0x40014000 in ?? ()
(gdb) continue
Continuing.

Breakpoint 1, jump_to_image_no_args (spl_image=0x80000000 )
at arch/arm/cpu/armv7/omap-common/boot-common.c:229
(gdb) si
0x80100000 in ?? ()
```

Bootloader instrumentation in action

(with GDB and QEMU)

```
(gdb) target remote | qemu-system-arm -gdb stdio -M beaglexm -sd sd.img -S
Remote debugging using | qemu-system-arm -gdb stdio -M beaglexm -sd sd.img -S
0x40014000 in ?? ()
(gdb) continue
Continuing.

Breakpoint 1, jump_to_image_no_args (spl_image=0x80000000 )
at arch/arm/cpu/armv7/omap-common/boot-common.c:229
(gdb) si
0x80100000 in ?? ()
(gdb) file u-boot
A program is being debugged already.
Are you sure you want to change the file? (y or n) y
Load new symbol table from "u-boot"? (y or n) y
Reading symbols from u-boot...done.
Error in re-setting breakpoint 1: Function "jump_to_image_no_args" not defined.
```

Bootloader instrumentation in action

(with GDB and QEMU)

```
(gdb) target remote | qemu-system-arm -gdb stdio -M beaglexm -sd sd.img -S
Remote debugging using | qemu-system-arm -gdb stdio -M beaglexm -sd sd.img -S
0x40014000 in ?? ()
(gdb) continue
Continuing.

Breakpoint 1, jump_to_image_no_args (spl_image=0x80000000 )
at arch/arm/cpu/armv7/omap-common/boot-common.c:229
(gdb) si
0x80100000 in ?? ()
(gdb) file u-boot
A program is being debugged already.
Are you sure you want to change the file? (y or n) y
Load new symbol table from "u-boot"? (y or n) y
Reading symbols from u-boot...done.
Error in re-setting breakpoint 1: Function "jump_to_image_no_args" not defined.
(gdb) break relocate_done
Breakpoint 2 at 0x801020b4: file arch/arm/lib/relocate.S, line 134.
```

Bootloader instrumentation in action

(with GDB and QEMU)

```
(gdb) target remote | qemu-system-arm -gdb stdio -M beaglexm -sd sd.img -S
Remote debugging using | qemu-system-arm -gdb stdio -M beaglexm -sd sd.img -S
0x40014000 in ?? ()
(gdb) continue
Continuing.

Breakpoint 1, jump_to_image_no_args (spl_image=0x80000000 )
at arch/arm/cpu/armv7/omap-common/boot-common.c:229
(gdb) si
0x80100000 in ?? ()
(gdb) file u-boot
A program is being debugged already.
Are you sure you want to change the file? (y or n) y
Load new symbol table from "u-boot"? (y or n) y
Reading symbols from u-boot...done.
Error in re-setting breakpoint 1: Function "jump_to_image_no_args" not defined.
(gdb) break relocate_done
Breakpoint 2 at 0x801020b4: file arch/arm/lib/relocate.S, line 134.
(gdb) continue
Continuing.

Breakpoint 2, relocate_done () at arch/arm/lib/relocate.S:134
134 in arch/arm/lib/relocate.S
```

Bootloader instrumentation in action

(with GDB and QEMU)

```
(gdb) target remote | qemu-system-arm -gdb stdio -M beaglexm -sd sd.img -S
Remote debugging using | qemu-system-arm -gdb stdio -M beaglexm -sd sd.img -S
0x40014000 in ?? ()
(gdb) continue
Continuing.

Breakpoint 1, jump_to_image_no_args (spl_image=0x80000000 )
at arch/arm/cpu/armv7/omap-common/boot-common.c:229
(gdb) si
0x80100000 in ?? ()
(gdb) file u-boot
A program is being debugged already.
Are you sure you want to change the file? (y or n) y
Load new symbol table from "u-boot"? (y or n) y
Reading symbols from u-boot...done.
Error in re-setting breakpoint 1: Function "jump_to_image_no_args" not defined.

(gdb) break relocate_done
Breakpoint 2 at 0x801020b4: file arch/arm/lib/relocate.S, line 134.
(gdb) continue
Continuing.

Breakpoint 2, relocate_done () at arch/arm/lib/relocate.S:134
134 in arch/arm/lib/relocate.S
(gdb) break board_init_r
Breakpoint 3 at 0x80104e7c: file common/board_r.c, line 957.
```

Bootloader instrumentation in action

(with GDB and QEMU)

```
(gdb) target remote | qemu-system-arm -gdb stdio -M beaglexm -sd sd.img -S
Remote debugging using | qemu-system-arm -gdb stdio -M beaglexm -sd sd.img -S
0x40014000 in ?? ()
(gdb) continue
Continuing.

Breakpoint 1, jump_to_image_no_args (spl_image=0x80000000 )
at arch/arm/cpu/armv7/omap-common/boot-common.c:229
(gdb) si
0x80100000 in ?? ()
(gdb) file u-boot
A program is being debugged already.
Are you sure you want to change the file? (y or n) y
Load new symbol table from "u-boot"? (y or n) y
Reading symbols from u-boot...done.
Error in re-setting breakpoint 1: Function "jump_to_image_no_args" not defined.

(gdb) break relocate_done
Breakpoint 2 at 0x801020b4: file arch/arm/lib/relocate.S, line 134.
(gdb) continue
Continuing.

Breakpoint 2, relocate_done () at arch/arm/lib/relocate.S:134
134 in arch/arm/lib/relocate.S
(gdb) break board_init_r
Breakpoint 3 at 0x80104e7c: file common/board_r.c, line 957.
(gdb) continue
Continuing.
```

Bootloader instrumentation in action

(with GDB and QEMU)

```
(gdb) target remote | qemu-system-arm -gdb stdio -M beaglexm -sd sd.img -S
Remote debugging using | qemu-system-arm -gdb stdio -M beaglexm -sd sd.img -S
0x40014000 in ?? ()
(gdb) continue
Continuing.

Breakpoint 1, jump_to_image_no_args (spl_image=0x80000000 )
at arch/arm/cpu/armv7/omap-common/boot-common.c:229
(gdb) si
0x80100000 in ?? ()
(gdb) file u-boot
A program is being debugged already.
Are you sure you want to change the file? (y or n) y
Load new symbol table from "u-boot"? (y or n) y
Reading symbols from u-boot...done.
Error in re-setting breakpoint 1: Function "jump_to_image_no_args" not defined.

(gdb) break relocate_done
Breakpoint 2 at 0x801020b4: file arch/arm/lib/relocate.S, line 134.
(gdb) continue
Continuing.

Breakpoint 2, relocate_done () at arch/arm/lib/relocate.S:134
134 in arch/arm/lib/relocate.S
(gdb) break board_init_r
Breakpoint 3 at 0x80104e7c: file common/board_r.c, line 957.
(gdb) continue
Continuing.

^Comap_gpmc_write: bad SDRAM idle mode 3
omap_i2c_write: Bad register 0x0000cc
```

How to debug a self-relocating bootloader

As according to U-Boot's doc/README.arm-relocation

start debugger

```
[hs@pollux u-boot]$ arm-linux-gdb u-boot
```

How to debug a self-relocating bootloader

As according to U-Boot's doc/README.arm-relocation

start debugger

```
[hs@pollux u-boot]$ arm-linux-gdb u-boot
```

connect to target

```
(gdb) target remote localhost:4444
```

How to debug a self-relocating bootloader

As according to U-Boot's doc/README.arm-relocation

start debugger

```
[hs@pollux u-boot]$ arm-linux-gdb u-boot
```

connect to target

```
(gdb) target remote localhost:4444
```

execute until relocation complete

```
(gdb) break _relocation_done  
(gdb) c
```

How to debug a self-relocating bootloader

As according to U-Boot's doc/README.arm-relocation

start debugger

```
[hs@pollux u-boot]$ arm-linux-gdb u-boot
```

connect to target

```
(gdb) target remote localhost:4444
```

execute until relocation complete

```
(gdb) break _relocation_done  
(gdb) c
```

discard symbol-file

```
(gdb) symbol-file  
Discard symbol table from `/home/hs/celf/u-boot/u-boot'? (y or n) y  
No symbol file now.
```

How to debug a self-relocating bootloader

As according to U-Boot's doc/README.arm-relocation

start debugger

```
[hs@pollux u-boot]$ arm-linux-gdb u-boot
```

connect to target

```
(gdb) target remote localhost:4444
```

execute until relocation complete

```
(gdb) break _relocation_done  
(gdb) c
```

discard symbol-file

```
(gdb) symbol-file  
Discard symbol table from `/home/hs/celf/u-boot/u-boot'? (y or n) y  
No symbol file now.
```

load new symbol table at relocated address

```
gdb) add-symbol-file u-boot 0x8ff08000  
add symbol table from file "u-boot" at  
  .text_addr = 0x8ff08000  
(y or n) y  
Reading symbols from /home/hs/celf/u-boot/u-boot...done.
```

How to debug a self-relocating bootloader

As according to U-Boot's doc/README.arm-relocation

start debugger

```
[hs@pollux u-boot]$ arm-linux-gdb u-boot
```

connect to target

```
(gdb) target remote localhost:1234
```

execute

```
(gdb) break _relocation_start
```

```
(gdb) c
```

(gdb) symbol-file u-boot.symbols

Discard symbol table file "u-boot.symbols".
No symbol file now.

load new symbols

```
(gdb) add-symbol-file u-boot.symbols
```

shutterstock.com - 296293568

```
add symbol table from file "u-boot" at  
.text_addr = 0xBff08000
```

```
(y or n) y
```

```
Reading symbols from /home/hs/celf/u-boot/u-boot...done.
```



address

"Demystifying" magic numbers

As suggested by the U-Boot developers

```
Program received signal SIGSTOP, Stopped (signal).
0x8ff17f18 in serial_getc () at serial_mxc.c:192
192     while (_REG(UART_PHYS + UTS) & UTS_RXEMPTY);
(gdb) add-symbol-file u-boot 0x8ff08000
```

"Demystifying" magic numbers

As suggested by the U-Boot developers

```
Program received signal SIGSTOP, Stopped (signal).
0x8ff17f18 in serial_getc () at serial_mxc.c:192
192     while (_REG(UART_PHYS + UTS) & UTS_RXEMPTY);
(gdb) add-symbol-file u-boot 0x8ff08000
```

"Demystifying" magic numbers

As suggested by the U-Boot developers

```
Program received signal SIGSTOP, Stopped (signal).
0x8ff17f18 in serial_getc () at serial_mxc.c:192
192     while (_REG(UART_PHYS + UTS) & UTS_RXEMPTY);
(gdb) add-symbol-file u-boot 0x8ff08000
```

"Get this address from u-boot bdinfo command or get it from gd->relocaddr in gdb. Interrupt execution by any means and re-load the symbols at the location specified by gd->relocaddr -- this is only valid after board_init_f."

"Demystifying" magic numbers

As suggested by the U-Boot developers

```
Program received signal SIGSTOP, Stopped (signal).
0x8ff17f18 in serial_getc () at serial_mxc.c:192
192     while (_REG(UART_PHYS + UTS) & UTS_RXEMPTY);
(gdb) add-symbol-file u-boot 0x8ff08000
```

"Get this address from u-boot bdinfo command or get it from gd->relocaddr in gdb. Interrupt execution by any means and re-load the symbols at the location specified by gd->relocaddr -- this is only valid after board_init_f."

```
=> bdinfo
arch_number = XXXXXXXXXX
boot_params = XXXXXXXXXX
DRAM bank   = XXXXXXXXXX
-> start    = XXXXXXXXXX
-> size     = XXXXXXXXXX
ethaddr     = XXXXXXXXXX
ip_addr     = XXXXXXXXXX
baudrate    = XXXXXXXXXX
TLB addr    = XXXXXXXXXX
relocaddr   = 0x8ff08000
reloc off   = XXXXXXXXXX
irq_sp      = XXXXXXXXXX
sp start    = XXXXXXXXXX
FB base     = XXXXXXXXXX
```

"Demystifying" magic numbers

As suggested by the U-Boot developers

```
Program received signal SIGSTOP, Stopped (signal).
0x8ff17f18 in serial_getc () at serial_mxc.c:192
192     while (_REG(UART_PHYS + UTS) & UTS_RXEMPTY);
(gdb) add-symbol-file u-boot 0x8ff08000
```

"Get this address from u-boot bdinfo command or get it from gd->relocaddr in gdb. Interrupt execution by any means and re-load the symbols at the location specified by gd->relocaddr -- this is only valid after board_init_f."

```
=> bdinfo
arch_number = XXXXXXXXXX
boot_params = XXXXXXXXXX
DRAM bank   = XXXXXXXXXX
-> start    = XXXXXXXXXX
-> size     = XXXXXXXXXX
ethaddr     = XXXXXXXXXX
ip_addr     = XXXXXXXXXX
baudrate    = XXXXXXXXXX
TLB addr    = XXXXXXXXXX
relocaddr   = 0x8ff08000
reloc off   = XXXXXXXXXX
irq_sp      = XXXXXXXXXX
sp start    = XXXXXXXXXX
FB base     = XXXXXXXXXX
```



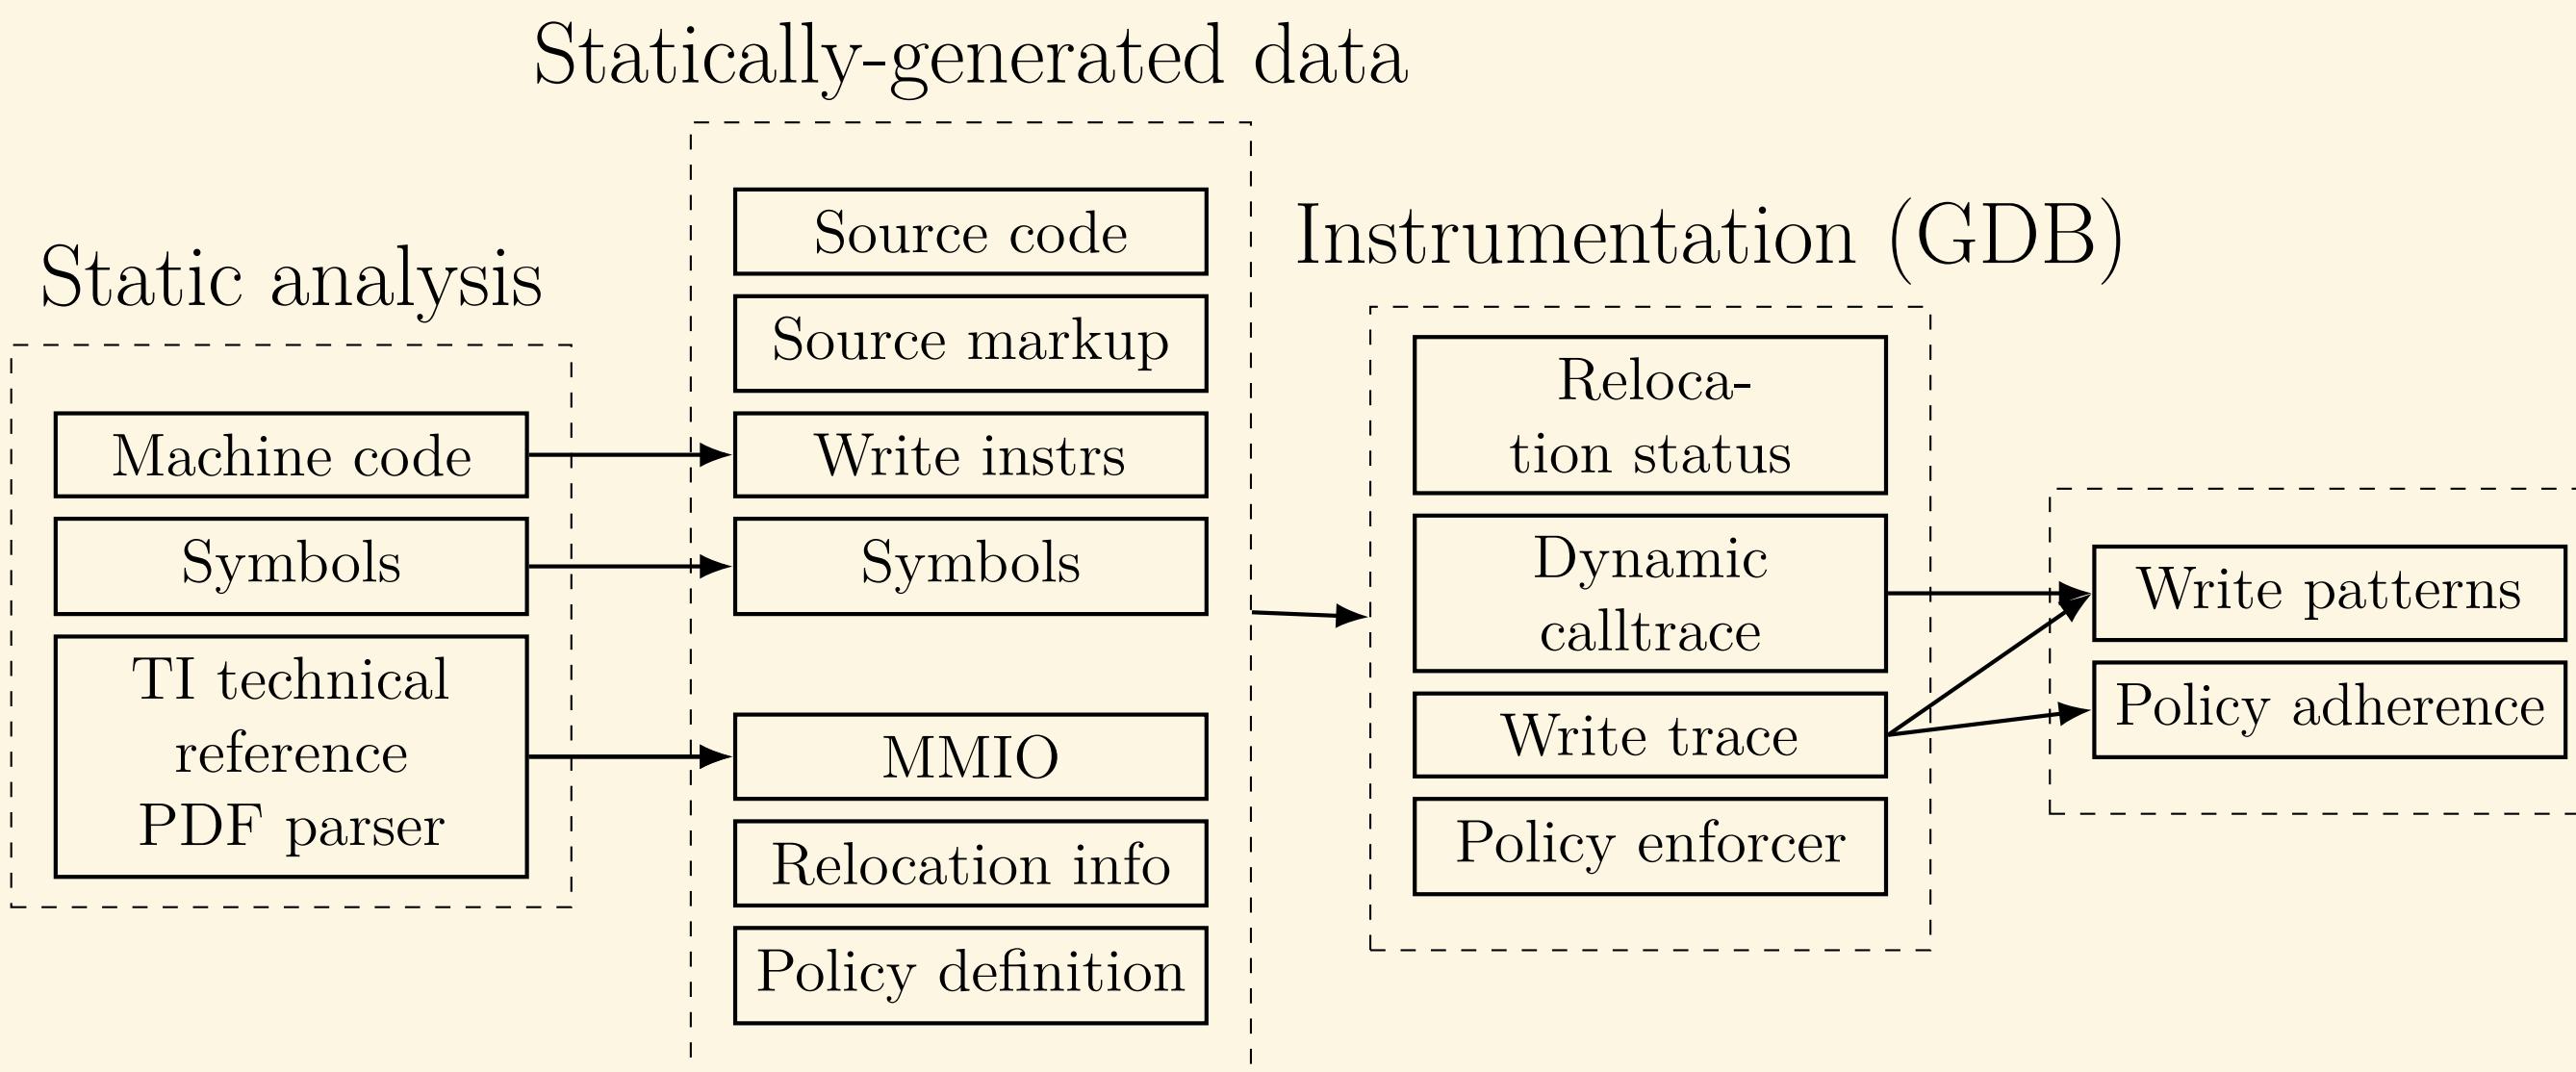
Can we do better?

The Fiddle toolsuite

Featuring:

- Static and dynamic analysis
- Instrumentation (via GDB)
 - Mediates **all** memory writes
- Language to express (and enforce) memory write policies
- Supports 32-bit ARM & amd64 (some support for arm64 & x86)
- Not just for bootloaders

MEDIATE ALL THE WRITES!



Source code at <https://typedregions.com>

Phase two

Relocation reconnaissance

Featuring the U-Boot SPL for BeagleBoard-xM

Tools at <https://typedregions.com>

Identifying relocation phases

Hypothesis:

The relocation process makes use of memcp-like operations

(Copying a sequence of adjacent bytes in a tight loop)

I call these **block write operations**

Block write operation

(\$ip, offset in image, destination, size, call stack)



Identifying relocation phases

Hypothesis:

The relocation process makes use of memcp-like operations

(Copying a sequence of adjacent bytes in a tight loop)

I call these **block write operations**

Block write operation

(\$ip, offset in image, destination, size, call stack)

Corollary:

We should be able to identify relocation phases from large block write operations



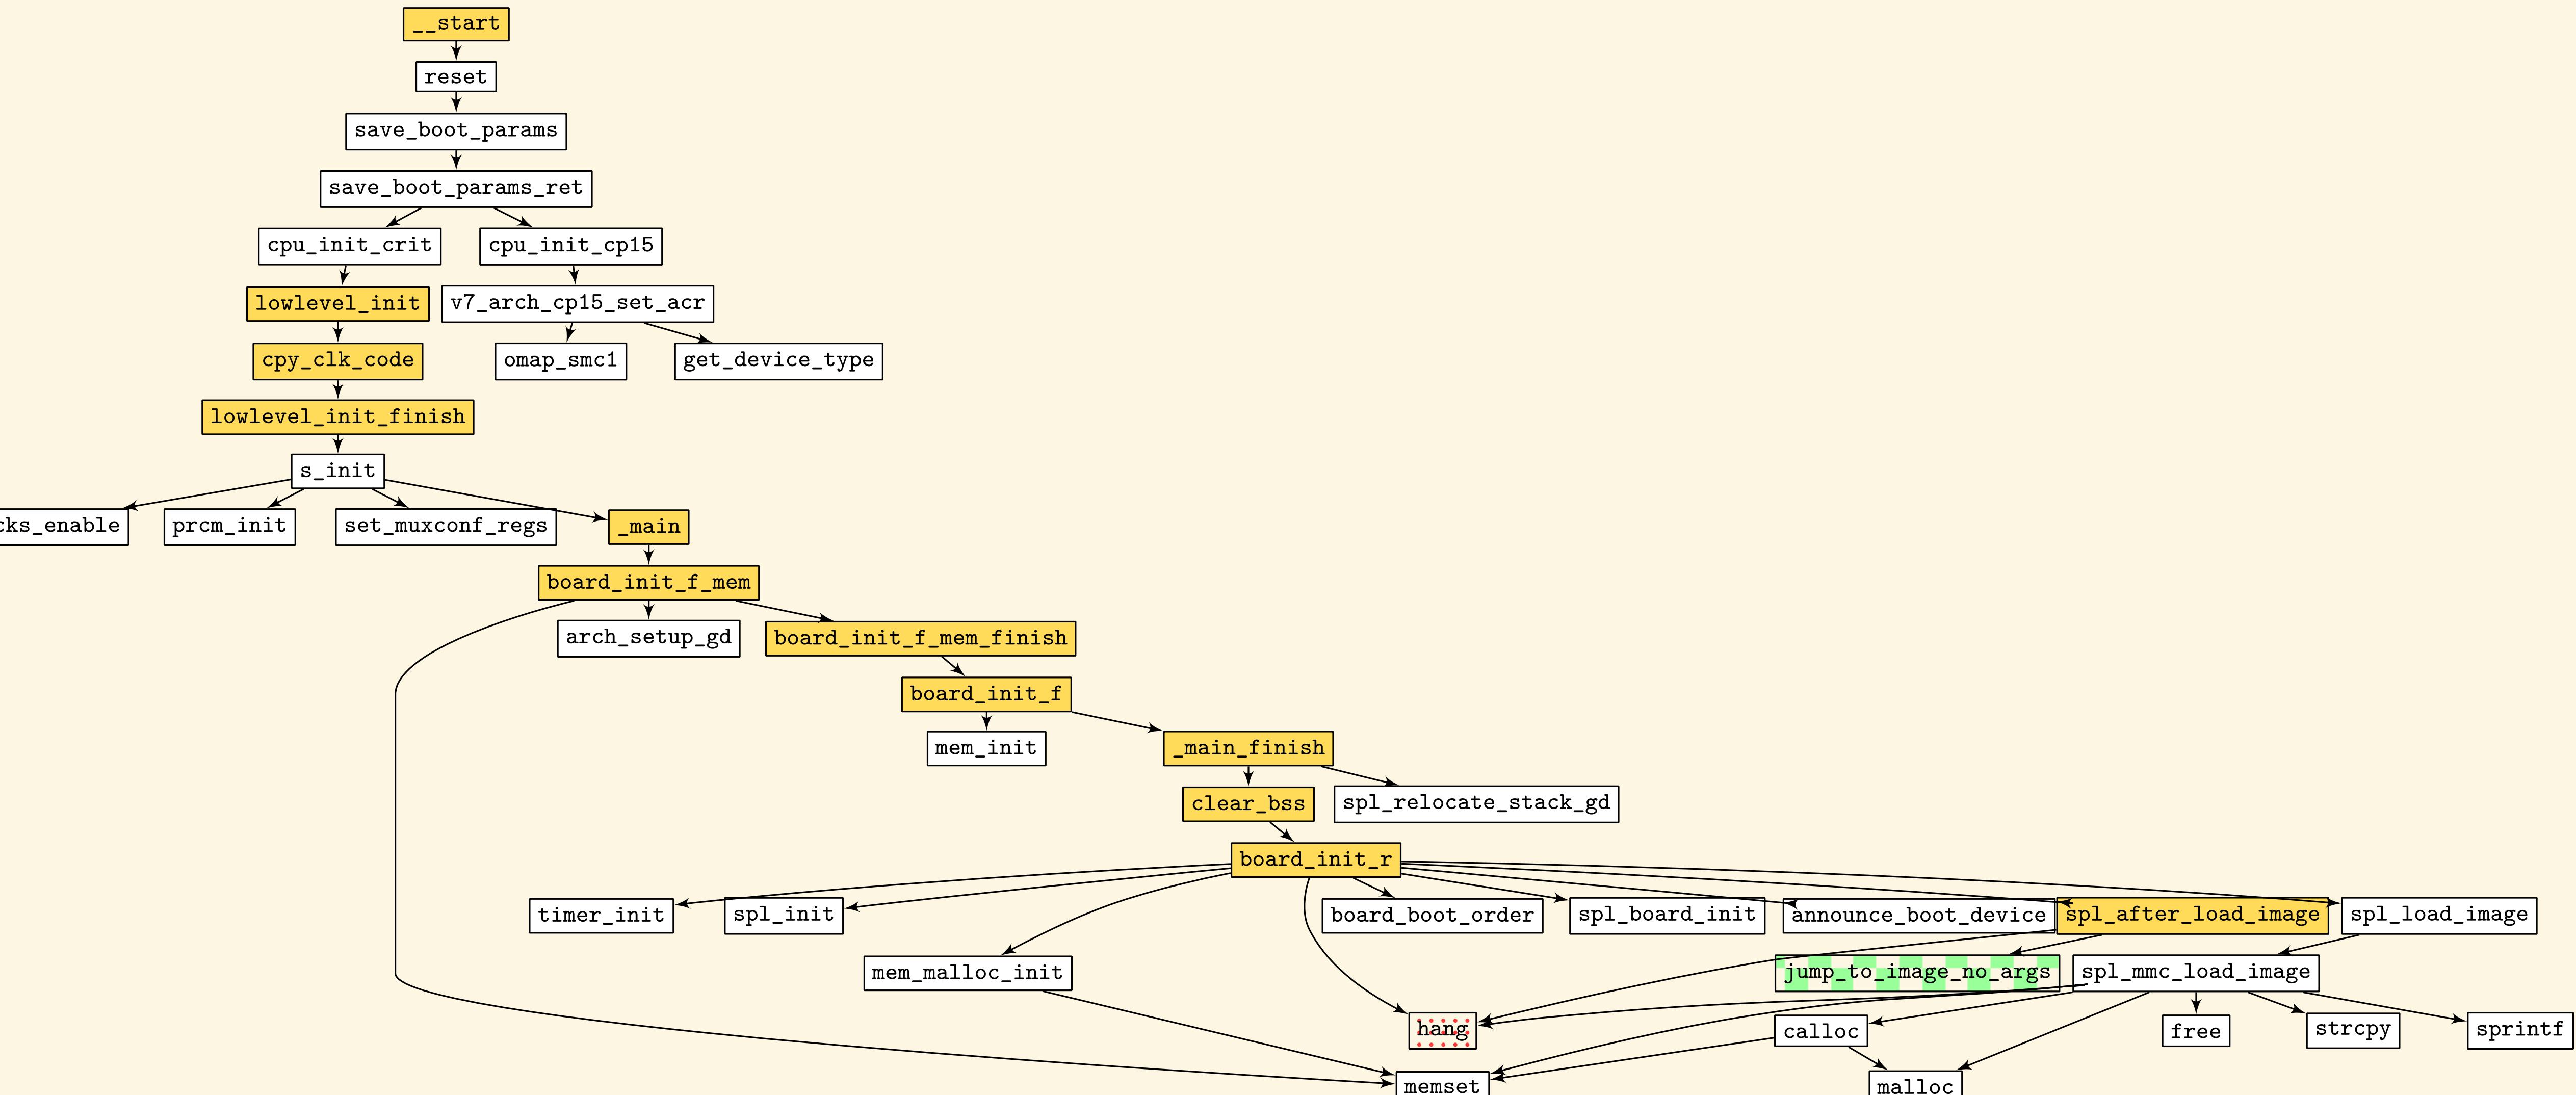
Block writes in U-Boot

```
[262144x] memset @ pc=0x40208a3e wrote 1048576 B to 0x80208000 str  
[49233x] clbss_l @ pc=0x402025c4 wrote 196932 B to 0x80000000 strc<-- Zero BSS  
[128x] mmc_read_data @ pc=0x40206cfa wrote 512 B to 0x4020f2c0  
str.w  
[128x] mmc_read_data @ pc=0x40206cfa wrote 512 B to 0x4020f2c0      <-- Read target image to memory  
str.w  
[128x] mmc_read_data @ pc=0x40206cfa wrote 512 B to 0x80104bc0  
str.w  
[128x] mmc_read_data @ pc=0x40206cfa wrote 512 B to 0x80104dc0  
str.w  
[83x] memset @ pc=0x40208a3e wrote 332 B to 0x80208038 str      <-- Relocate "go_to_speed" function  
[9x] next2 @ pc=0x402009c4 wrote 288 B to 0x4020f840 stmia  
[54x] memset @ pc=0x40208a3e wrote 216 B to 0x4020fe10 str  
[30x] memcpy @ pc=0x40208a72 wrote 120 B to 0x800200c0 str  
[16x] memcpy @ pc=0x40208a72 wrote 64 B to 0x4020f118 str  
[1x] omap_smci @ pc=0x40200970 wrote 40 B to 0x40200234 push      <-- Relocate bookkeeping data  
                                         <-- Function call/stack push
```

# write operations	~400,000
# block writes	10,000

U-Boot's static call graph (for reference)

Generated using IDA Pro (and then simplified by hand)

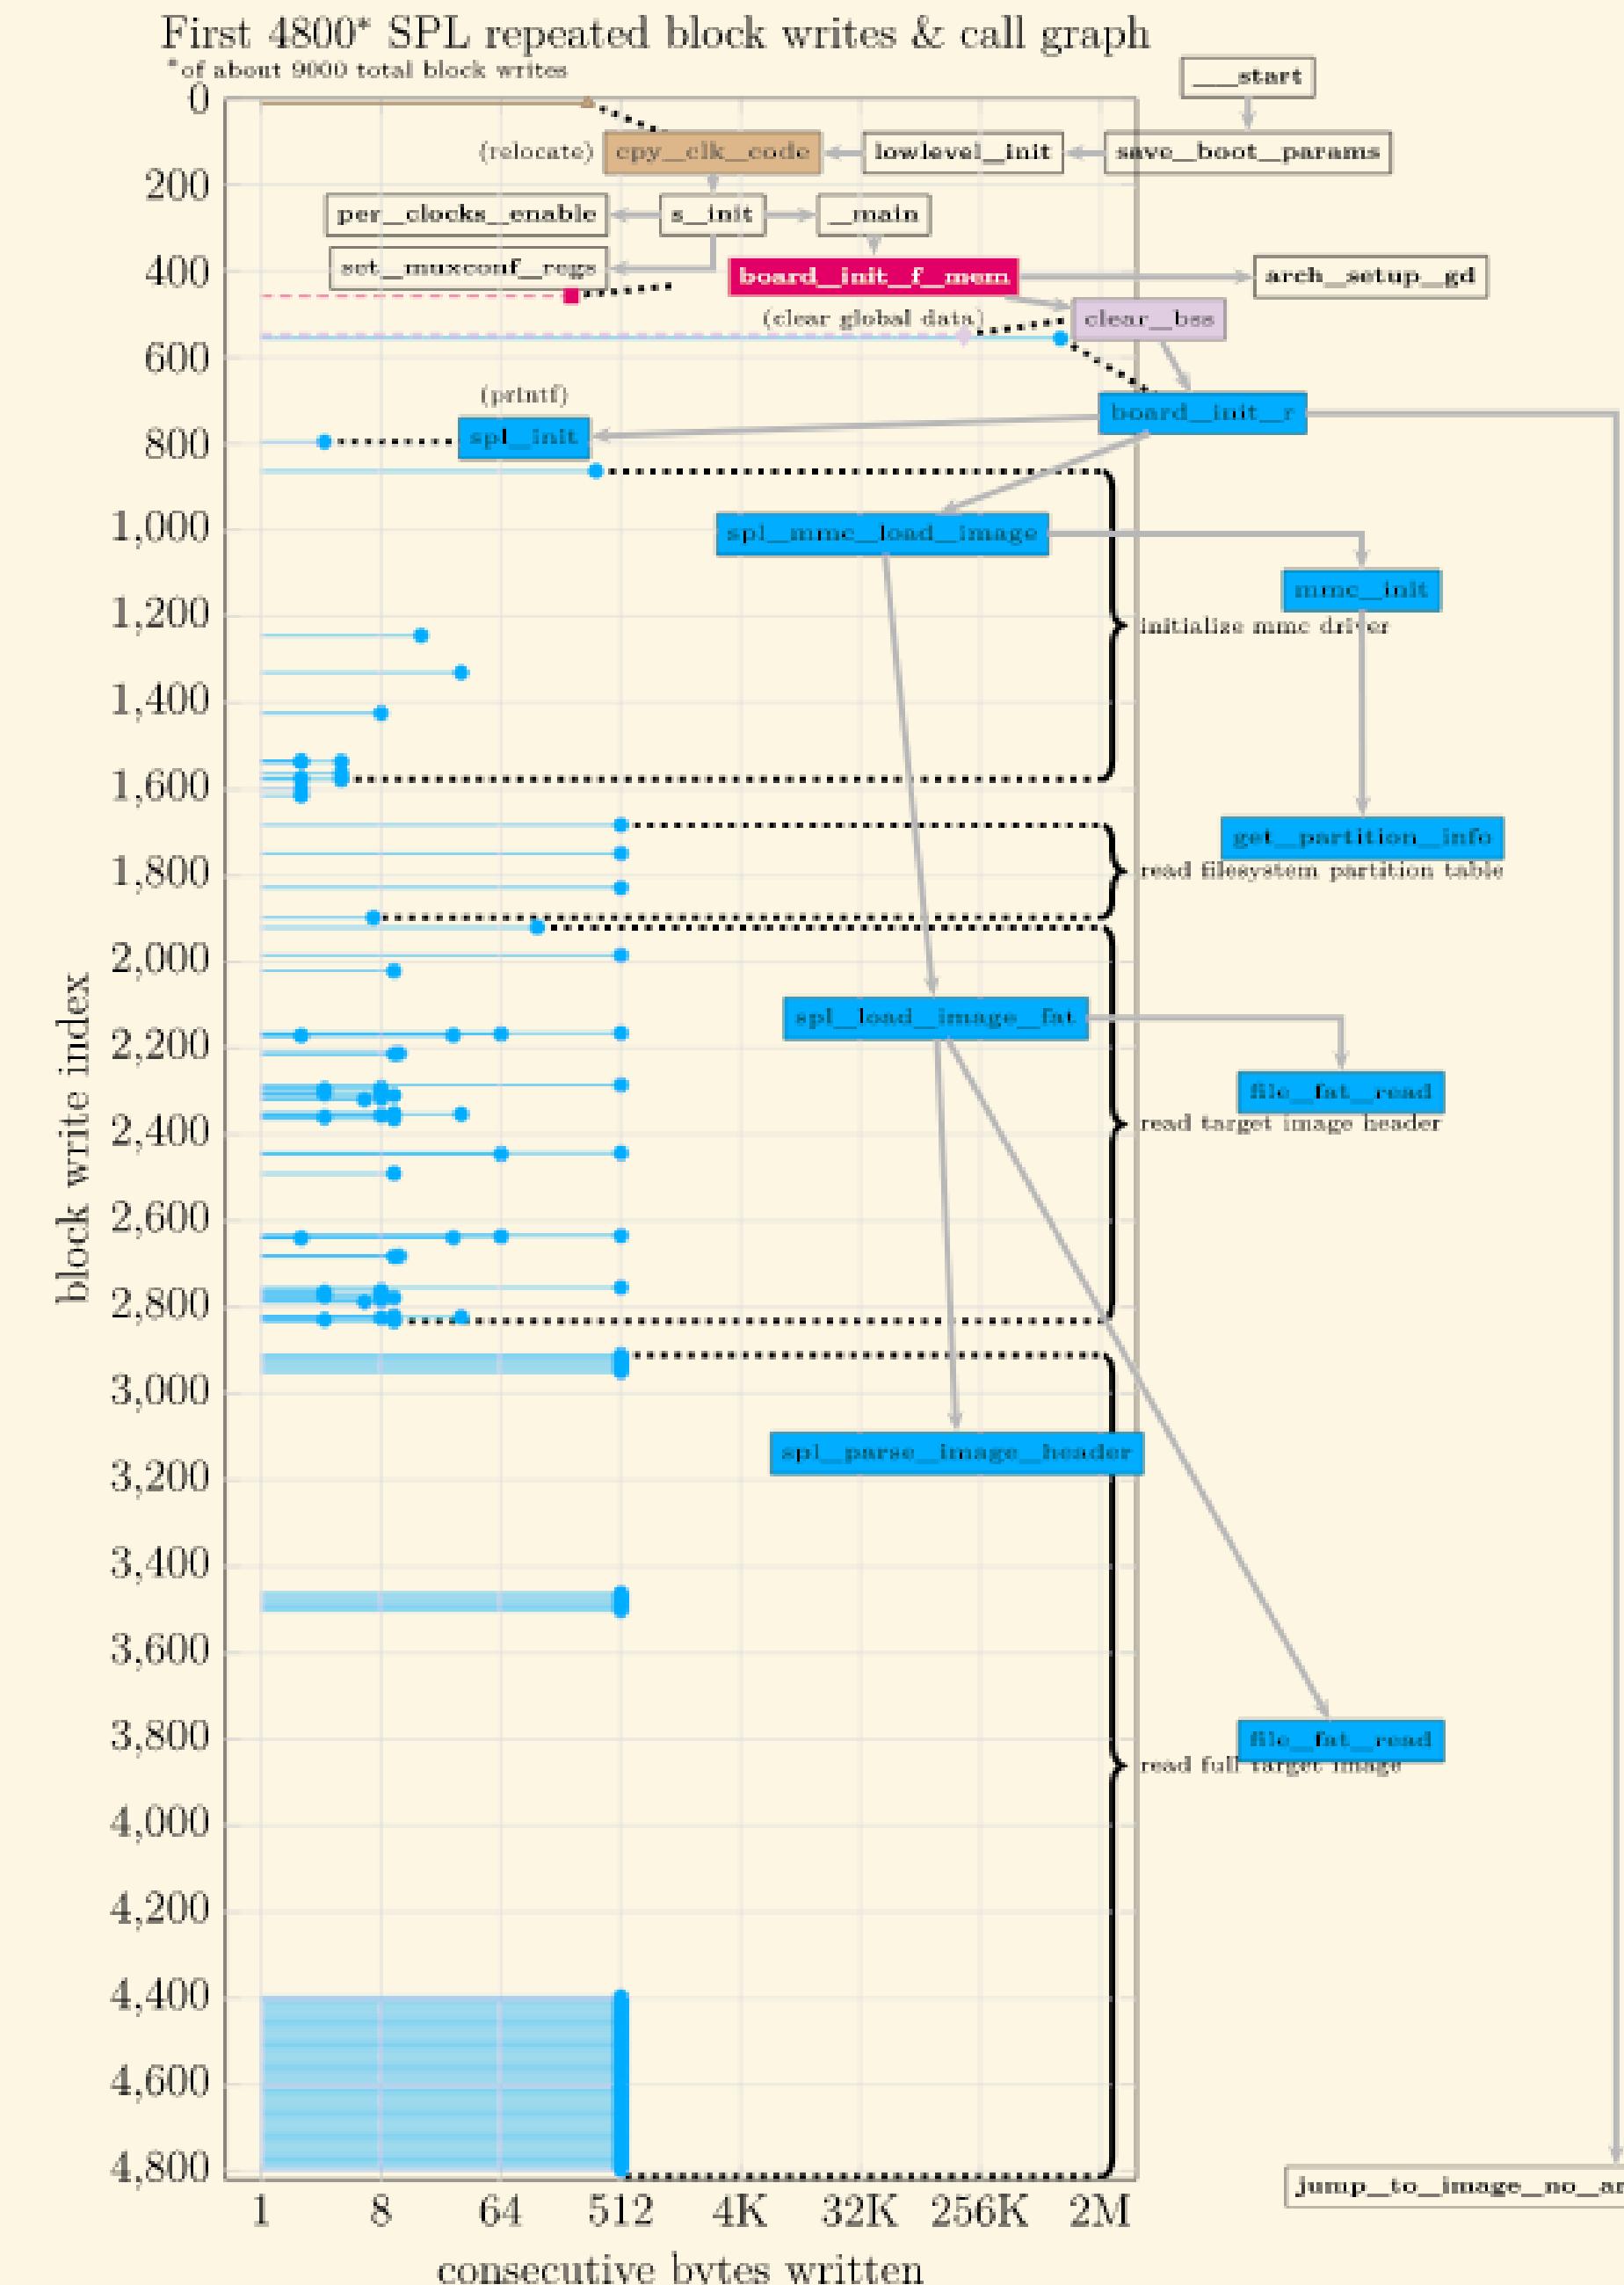


Calltrace of successful U-Boot execution

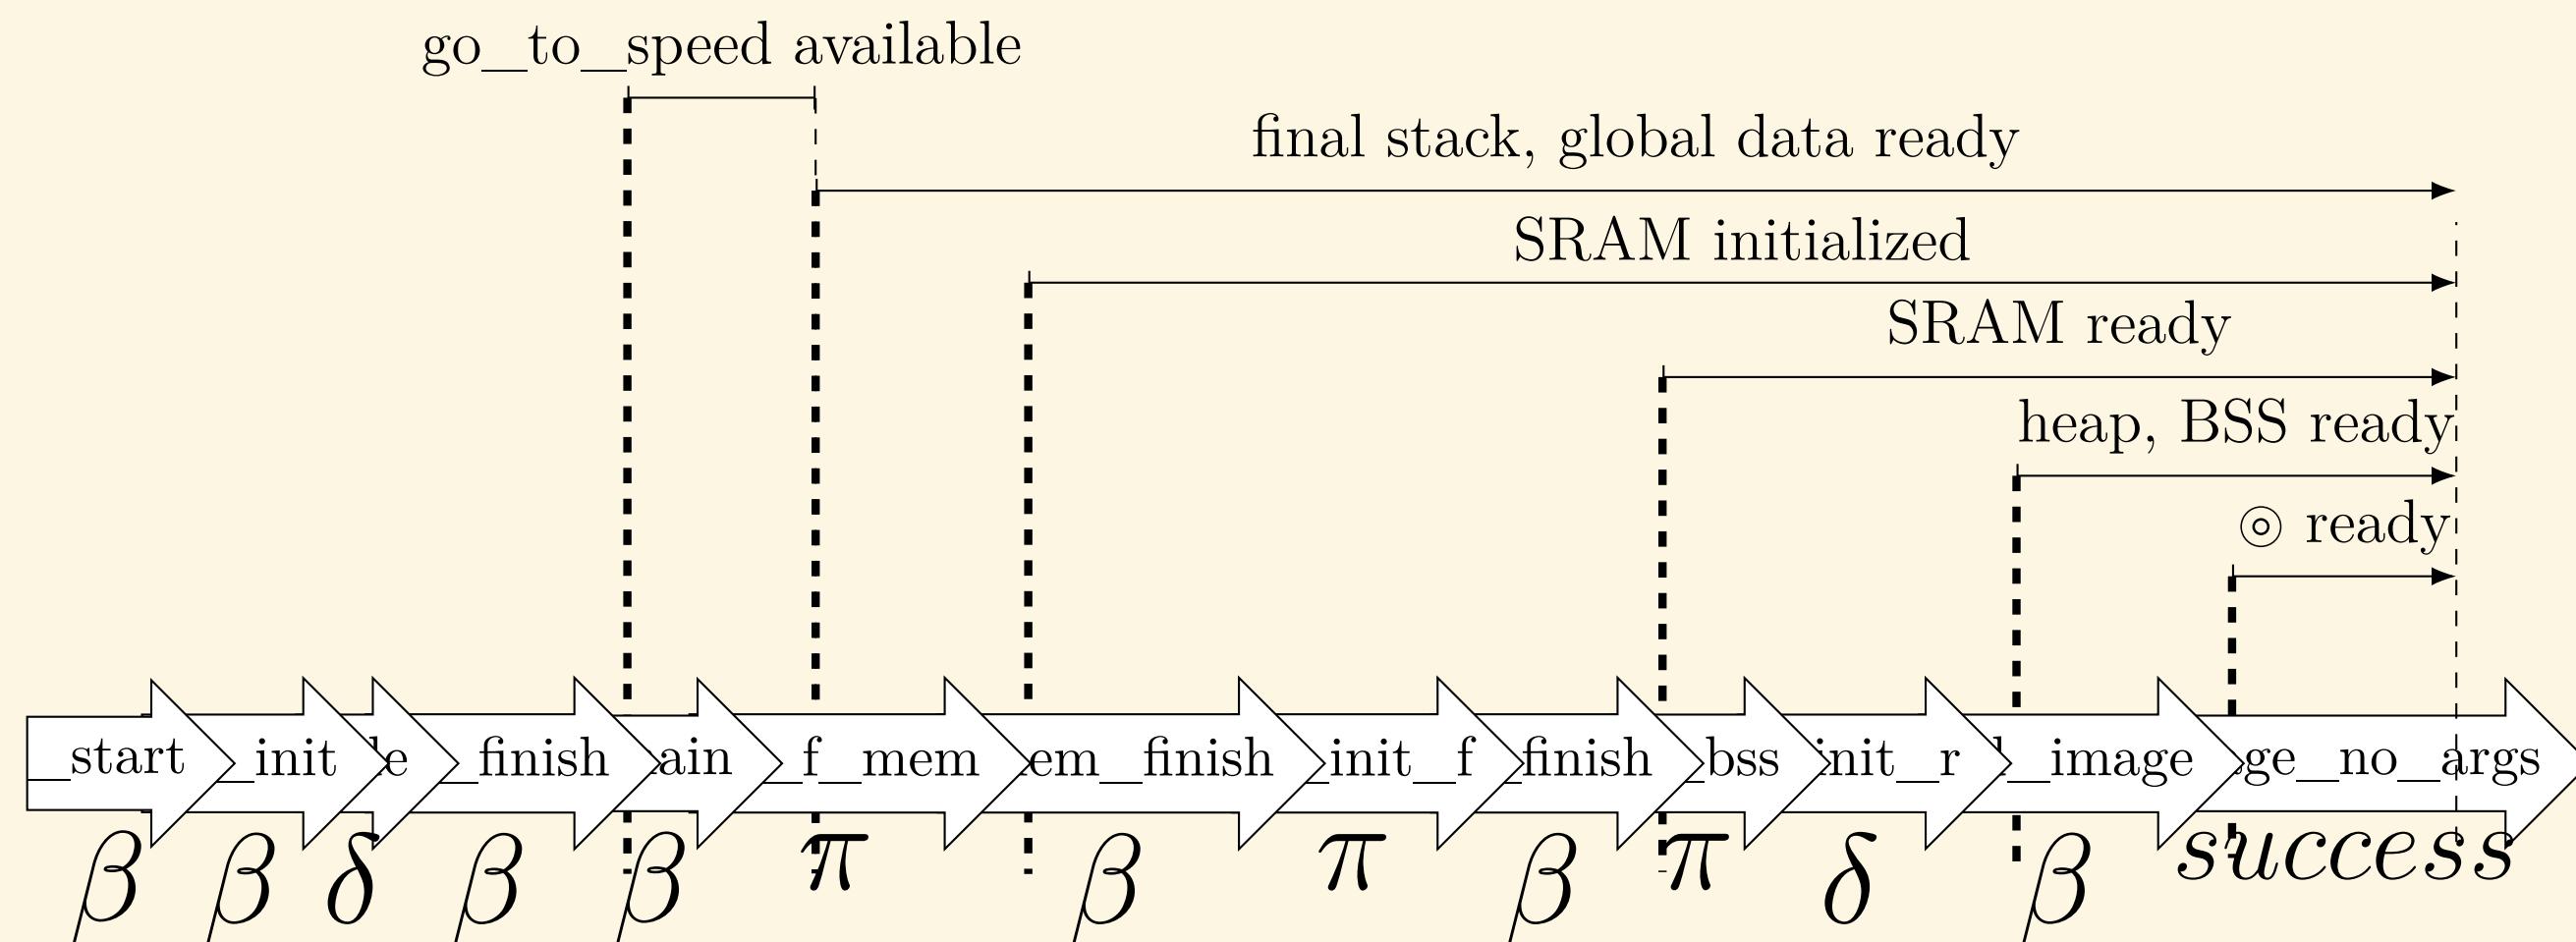
Example output from U-Boot SPL execution

```
• > save_boot_params {arch/ahrm/cpu/armv7/omap-common/lowlevel_init.S::22}
  ◦ > cpu_init_cp15 {arch/arm/cpu/armv7/start.S::120}
    ▪ > v7_arch_cp15_set_acr {arch/arm/cpu/armv7/omap3/board.c::388}
      ▪ > get_device_type {arch/arm/cpu/armv7/omap3/sys_info.c::247}
      ▪ < get_device_type
      ▪ > omap_smcl {arch/arm/cpu/armv7/omap-common/lowlevel_init.S::31}
      ▪ < omap_smcl
    ▪ < v7_arch_cp15_set_acr
    ▪ > v7_arch_cp15_set_acr {arch/arm/cpu/armv7/omap3/board.c::388}
      ▪ > get_device_type {arch/arm/cpu/armv7/omap3/sys_info.c::247}
      ▪ < get_device_type
      ▪ > omap_smcl {arch/arm/cpu/armv7/omap-common/lowlevel_init.S::31}
      ▪ < omap_smcl
    ▪ < v7_arch_cp15_set_acr
    ▪ > v7_arch_cp15_set_acr {arch/arm/cpu/armv7/omap3/board.c::388}
      ▪ > get_device_type {arch/arm/cpu/armv7/omap3/sys_info.c::247}
      ▪ < get_device_type
      ▪ > omap_smcl {arch/arm/cpu/armv7/omap-common/lowlevel_init.S::31}
      ▪ < omap_smcl
    ▪ < v7_arch_cp15_set_acr
  ◦ < cpu_init_cp15
  ◦ > cpu_init_crit {arch/arm/cpu/armv7/start.S::273}
    ▪ > lowlevel_init {arch/arm/cpu/armv7/omap3/lowlevel_init.S::187}
      ▪ > cpy_clk_code {arch/arm/cpu/armv7/omap3/lowlevel_init.S::49}
        ▪ > lowlevel_init_finish {arch/arm/cpu/armv7/omap3/lowlevel_init.S::203}
          ▪ > s_init {arch/arm/cpu/armv7/omap3/board.c::190}
            ▪ > watchdog_init {arch/arm/cpu/armv7/omap3/board.c::256}
              ▪ > wait_on_value {arch/arm/cpu/armv7/syslib.c::37}
              ▪ < wait_on_value
              ▪ > wait_for_command_complete {arch/arm/cpu/armv7/omap3/board.c::247}
              ▪ < wait_for_command_complete
            ▪ < watchdog_init
          ▪ > try_unlock_memory {arch/arm/cpu/armv7/omap3/board.c::152}
            ▪ > is_running_in_sdram {arch/arm/cpu/armv7/omap3/sys_info.c::226}
              ▪ > get_base {arch/arm/cpu/armv7/omap3/sys_info.c::191}
              ▪ < get_base
            ▪ < is_running_in_sdram
            ▪ > get_device_type {arch/arm/cpu/armv7/omap3/sys_info.c::247}
            ▪ < get_device_type
          ▪ > secure_unlock_mem {arch/arm/cpu/armv7/omap3/board.c::91}
```

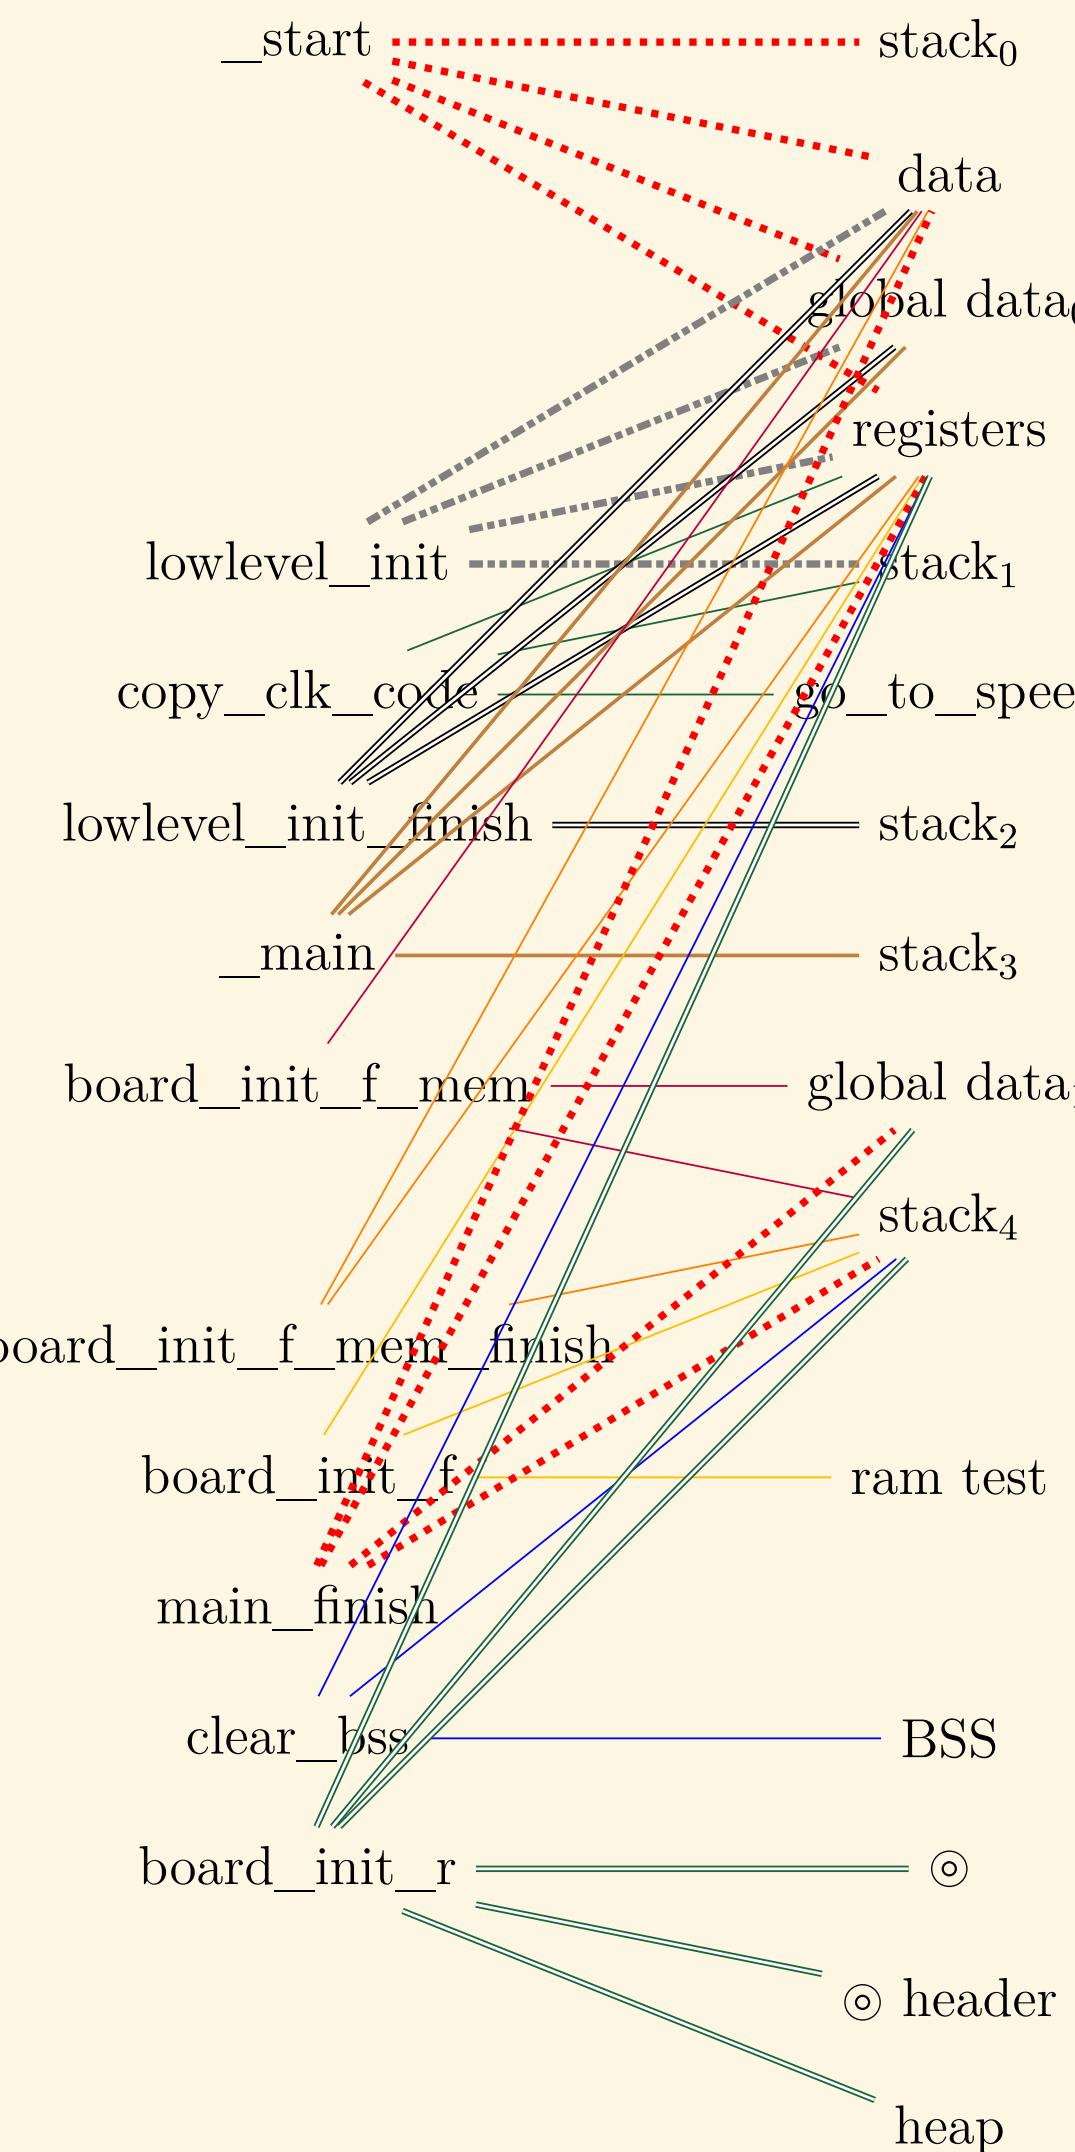
Calltrace and write data combined



Understanding U-Boot SPL execution as a sequence of phases



Devising code and data relationships



Other magic numbers

```
switch(beagle_revision()) {  
    case REVISION_C4:  
        if (identify_xm_ddr() == NUMONYX_MCP) {  
            /* 512MB/bank */  
            __raw_writel(0x4, SDRC_CS_CFG);  
  
            __raw_writel(SDP_SDRC_MDCFG_0_DDR_NUMONYX_XM,  
                         SDRC_MCFG_0);
```

```
argument = 0x0000 << 16;  
err = mmc_send_cmd(MMC_CMD55, argument, resp);  
if (err == 1) {  
    mmc_card_cur->card_type = SD_CARD;  
} else {  
    mmc_card_cur->card_type = MMC_CARD;
```

Other magic numbers

```
switch(beagle_revision()) {
    case REVISION_C4:
        if (identify_xm_ddr() == NUMONYX_MCP) {
            /* 512MB/bank */
            __raw_writel(0x4, SDRC_CS_CFG);
            __raw_writel(SDP_SDRC_MDCFG_0_DDR_NUMONYX_XM,
                         SDRC_MCFG_0);
```

```
argument = 0x0000 << 16;
err = mmc_send_cmd(MMC_CMD55, argument, resp);
if (err == 1) {
    mmc_card_cur->card_type = SD_CARD;
} else {
    mmc_card_cur->card_type = MMC_CARD;
```

So I built a PDF scraper

```
for o in page:
    if cls.is_not_in_table_bounds(tablestart, tableend, o):
        continue
    if isinstance(o, layout.LTRect) and (abs(o.bbox[1] - o.bbox[3]) < 1.5): # is a line
        # if left edge of line is not near left edge of name col_info
        if abs(namecol.bbox[0] - (o.bbox[0] + offset)) > 5:
            continue
        odiff = cls.vertical_offset(namecol, o)
        if (odiff > 0) and (closest_rect is None):
            closest_rect = o
```

Other magic numbers

```
switch(beagle_revision()) {
    case REVISION_C4:
        if (identify_xm_ddr() == NUMONYX_MCP) {
            /* 512MB/bank */
            __raw_writel(0x4, SDRC_CS_CFG);
            __raw_writel(SDP_SDRC_MDCFG_0_DDR_NUMONYX_XM,
                         SDRC_MCFG_0);
```

```
argument = 0x0000 << 16;
err = mmc_send_cmd(MMC_CMD55, argument, resp);
if (err == 1) {
    mmc_card_cur->card_type = SD_CARD;
} else {
    mmc_card_cur->card_type = MMC_CARD;
```

So I built a PDF scraper

```
for o in page:
    if cls.is_not_in_table_bounds(tablestart, tableend, o):
        continue
    if isinstance(o, layout.LTRect) and (abs(o.bbox[1] - o.bbox[3]) < 1.5): # is a line
        # if left edge of line is not near left edge of name col_info
        if abs(namecol.bbox[0] - (o.bbox[0] + offset)) > 5:
            continue
        odiff = cls.vertical_offset(namecol, o)
        if (odiff > 0) and (closest_rect is None):
            closest_rect = o
```

that transforms

Table 10-28. GPMC Registers Mapping Summary

Register Name	Type	Register Width (Bits)	Address Offset	Physical Address
GPMC_REVISION	R	32	0x0000 0000	0xE00 0000
GPMC_SYSCONFIG	RW	32	0x0000 0010	0xE00 0010
GPMC_SYSSTATUS	R	32	0x0000 0014	0xE00 0014
GPMC_IRQSTATUS	RW	32	0x0000 0018	0xE00 0018

Other magic numbers

```
switch(beagle_revision()) {
    case REVISION_C4:
        if (identify_xm_ddr() == NUMONYX_MCP) {
            /* 512MB/bank */
            __raw_writel(0x4, SDRC_CS_CFG);
            __raw_writel(SDP_SDRC_MDCFG_0_DDR_NUMONYX_XM,
                         SDRC_MCFG_0);
```

```
argument = 0x0000 << 16;
err = mmc_send_cmd(MMC_CMD55, argument, resp);
if (err == 1) {
    mmc_card_cur->card_type = SD_CARD;
} else {
    mmc_card_cur->card_type = MMC_CARD;
```

So I built a PDF scraper that transforms

```
for o in page:
    if cls.is_not_in_table_bounds(tablestart, tableend, o):
        continue
    if isinstance(o, layout.LTRect) and (abs(o.bbox[1] - o.bbox[3]) < 1.5): # is a line
        # if left edge of line is not near left edge of name col_info
        if abs(namecol.bbox[0] - (o.bbox[0] + offset)) > 5:
            continue
        odiff = cls.vertical_offset(namecol, o)
        if (odiff > 0) and (closest_rect is None):
            closest_rect = o
```

Table 10-28. GPMC Registers Mapping Summary

Register Name	Type	Register Width (Bits)	Address Offset	Physical Address
GPMC_REVISION	R	32	0x0000 0000	0xE00 0000
GPMC_SYSCONFIG	RW	32	0x0000 0010	0xE00 0010
GPMC_SYSSTATUS	R	32	0x0000 0014	0xE00 0014
GPMC_IRQSTATUS	RW	32	0x0000 0018	0xE00 0018

into

```
CM_FCLKEN_IVA2,RW,W,32,0x00000000,0x48004000,Table 3-93. IVA2_CM Register Summary
CM_CLKEN_PLL_IVA2,RW,W,32,0x00000004,0x48004004,Table 3-93. IVA2_CM Register Summary
CM_IDLEST_IVA2,R,C,32,0x00000020,0x48004020,Table 3-93. IVA2_CM Register Summary
CM_IDLEST_PLL_IVA2,R,C,32,0x00000024,0x48004024,Table 3-93. IVA2_CM Register Summary
CM_AUTOIDLE_PLL_IVA2,RW,W,32,0x00000034,0x48004034,Table 3-93. IVA2_CM Register Summary
```

A simpler example



Hello, world!

```
#include <stdio.h>
#include <string.h>

#define SIZE 512
char memory[SIZE];

void do_nothing() {}

void say_hello() {
    printf("Hello, world\n");
}

void modify_memory() {
    for (int i = 0; i <= SIZE; i++) {
        memory[i] = 'A';
    }
}

int main(int argc, char *argv[]) {
    say_hello();
    modify_memory();
    do_nothing();
    return 0;
}
```

(Built with -static,-no-pie)

Calltrace of Hello, world!

- > _libcstartmain {arm-linux-gnueabihf-glibc/src/glibc-2.27/csu/libc-start.c::137}
 - > dlauxinit {arm-linux-gnueabihf-glibc/src/glibc-2.27/elf/dl-support.c::226}
 - < dlauxinit
 - > _libcinitsecure {arm-linux-gnueabihf-glibc/src/glibc-2.27/elf/enbl-secure.c::32}
 - < _libcinitsecure
 - > _tunablesinit {arm-linux-gnueabihf-glibc/src/glibc-2.27/elf/dl-tunables.c::289}
 - < _tunablesinit
 - > _libcsetuptls {arm-linux-gnueabihf-glibc/src/glibc-2.27/csu/../csu/libc-tls.c::107}
 - > _udivsi3
 - > aeabiuidiv
 - < _udivsi3
 - < aeabiuidiv
 - > sbrk {arm-linux-gnueabihf-glibc/src/glibc-2.27/misc/sbrk.c::32}
 - > _sbrk {arm-linux-gnueabihf-glibc/src/glibc-2.27/misc/sbrk.c::32}
 - > _brk {arm-linux-gnueabihf-glibc/src/glibc-2.27/misc/..../sysdeps/unix/sysv/linux/arm/brk.c::28}
 - > brk {arm-linux-gnueabihf-glibc/src/glibc-2.27/misc/..../sysdeps/unix/sysv/linux/arm/brk.c::28}
 - < _brk
 - < brk
 - > _brk {arm-linux-gnueabihf-glibc/src/glibc-2.27/misc/..../sysdeps/unix/sysv/linux/arm/brk.c::28}
 - > brk {arm-linux-gnueabihf-glibc/src/glibc-2.27/misc/..../sysdeps/unix/sysv/linux/arm/brk.c::28}
 - < _brk
 - < brk
 - < sbrk
 - < _sbrk
 - > memcpy {arm-linux-gnueabihf-glibc/src/glibc-2.27/string/..../sysdeps/arm/memcpy.S::63}
 - < memcpy
 - > _udivsi3
 - > aeabiuidiv
 - < _udivsi3
 - < aeabiuidiv
 - < _libcsetuptls
 - > dldiscoverosversion {arm-linux-gnueabihf-glibc/src/glibc-2.27/elf/..../sysdeps/unix/sysv/linux/dl-sysdep.c::45}
 - > uname {arm-linux-gnueabihf-glibc/src/glibc-2.27 posix/..../sysdeps/unix/syscall-template.S::78}
 - > _uname {arm-linux-gnueabihf-glibc/src/glibc-2.27 posix/..../sysdeps/unix/syscall-template.S::78}
 - < uname
 - < _uname
 - < dldiscoverosversion
 - > _libcinitfirst {arm-linux-gnueabihf-glibc/src/glibc-2.27/csu/..../csu/init-first.c::44}
 - > dlondynamicinit {arm-linux-gnueabihf-glibc/src/glibc-2.27/elf/dl-support.c::309}
 - > dlgetorigin {arm-linux-gnueabihf-glibc/src/glibc-2.27/elf/..../sysdeps/unix/sysv/linux/dl-origin.c::36}

This is a javascript-enabled demo

Checking for unexpected writes

1. Configure instrumentation suite to work with sample
 2. Run sample through instrumentation suite's **static analysis**
 3. Construct policy to target "substages" and memory regions of interest
 4. Import policy
 5. Execute **dynamic analysis**
 6. Use **post-analysis** to highlight policy violations
-

Checking for unexpected writes

1. Configure instrumentation suite to work with sample
2. Run sample through instrumentation suite's **static analysis**
3. Construct policy to target "substages" and memory regions of interest
4. Import policy
5. Execute **dynamic analysis**
6. Use **post-analysis** to highlight policy violations

Region definitions

```
regions:  
ALL:  
  type: "global"  
  addresses: [0, 0xFFFFFFFF]  
  subregions:  
    buffer:  
      type: "global"  
      addresses: [0x8bb98, 0x8bd98]  
    stack:  
      type: "stack"  
      addresses: [0xffffe0000, 0xfffff0000]  
    ro:  
      type: "global"  
      addresses: [[0x0, 0x8bb898],  
                  [0x8bd98, 0xffffe0000],  
                  [0xfffff0000, 0xffffffff]]  
  
stagename: "_single"
```

Substage/policy definitions

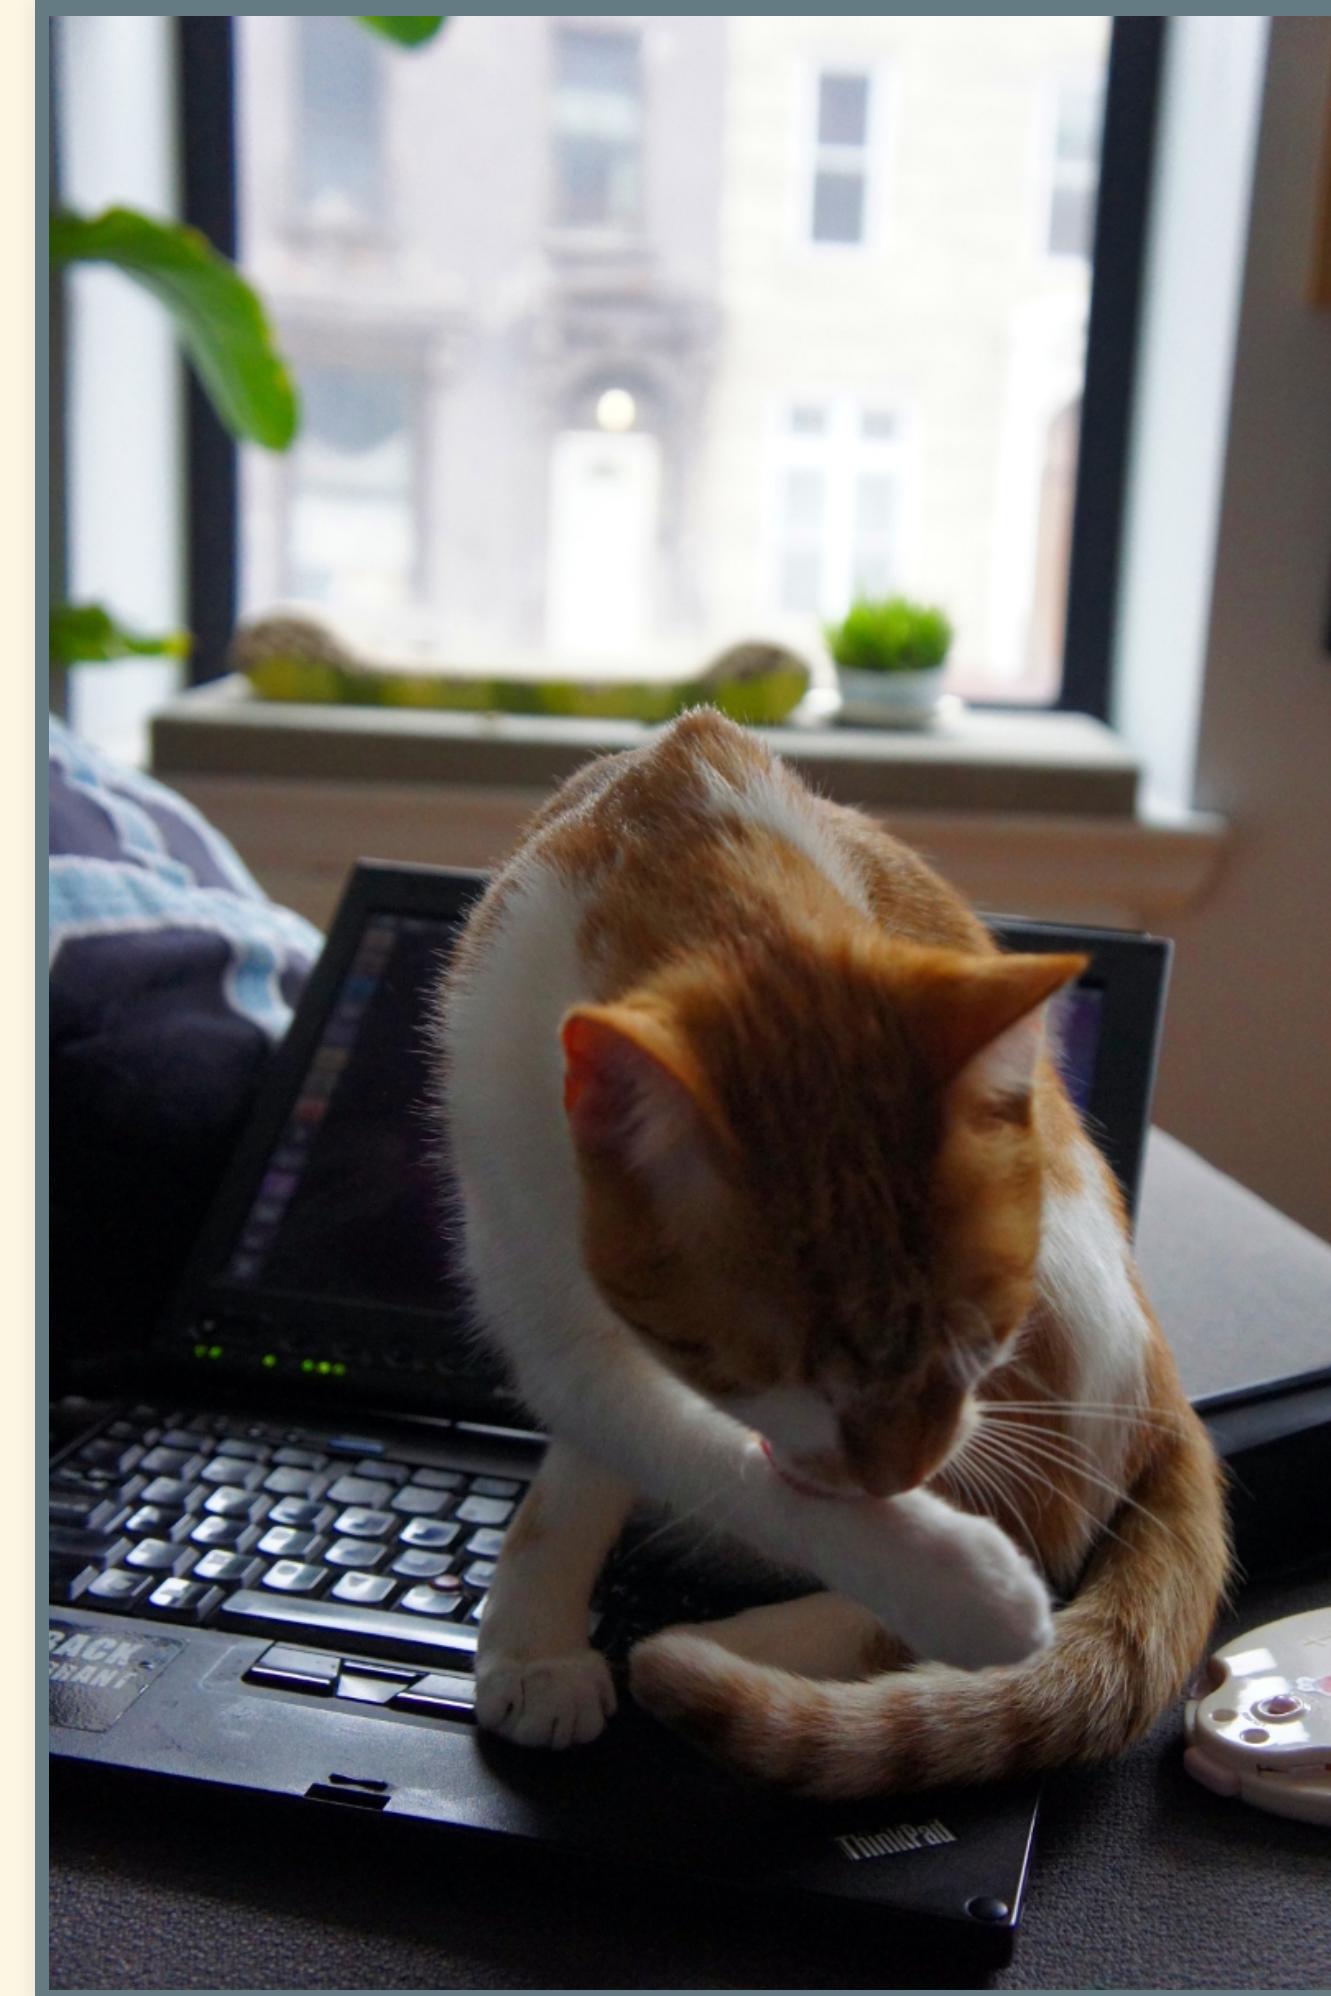
```
_start:  
  substage_type: "bookkeeping"  
  new_regions: ["ALL.buffer", "ALL.stack",  
"ALL.ro"]  
  reclassified_regions:  
    ALL.ro: "global"  
    ALL.stack: "stack"  
    ALL.buffer: "global"  
  
main:  
  substage_type: "bookkeeping"  
  
modify_memory:  
  substage_type: "bookkeeping"  
  reclassified_regions:  
    ALL.ro: "readonly"  
  
do_nothing:  
  substage_type: "bookkeeping"  
  reclassified_regions:  
    ALL.ro: "global"
```

Phase three

An attempt at a demo

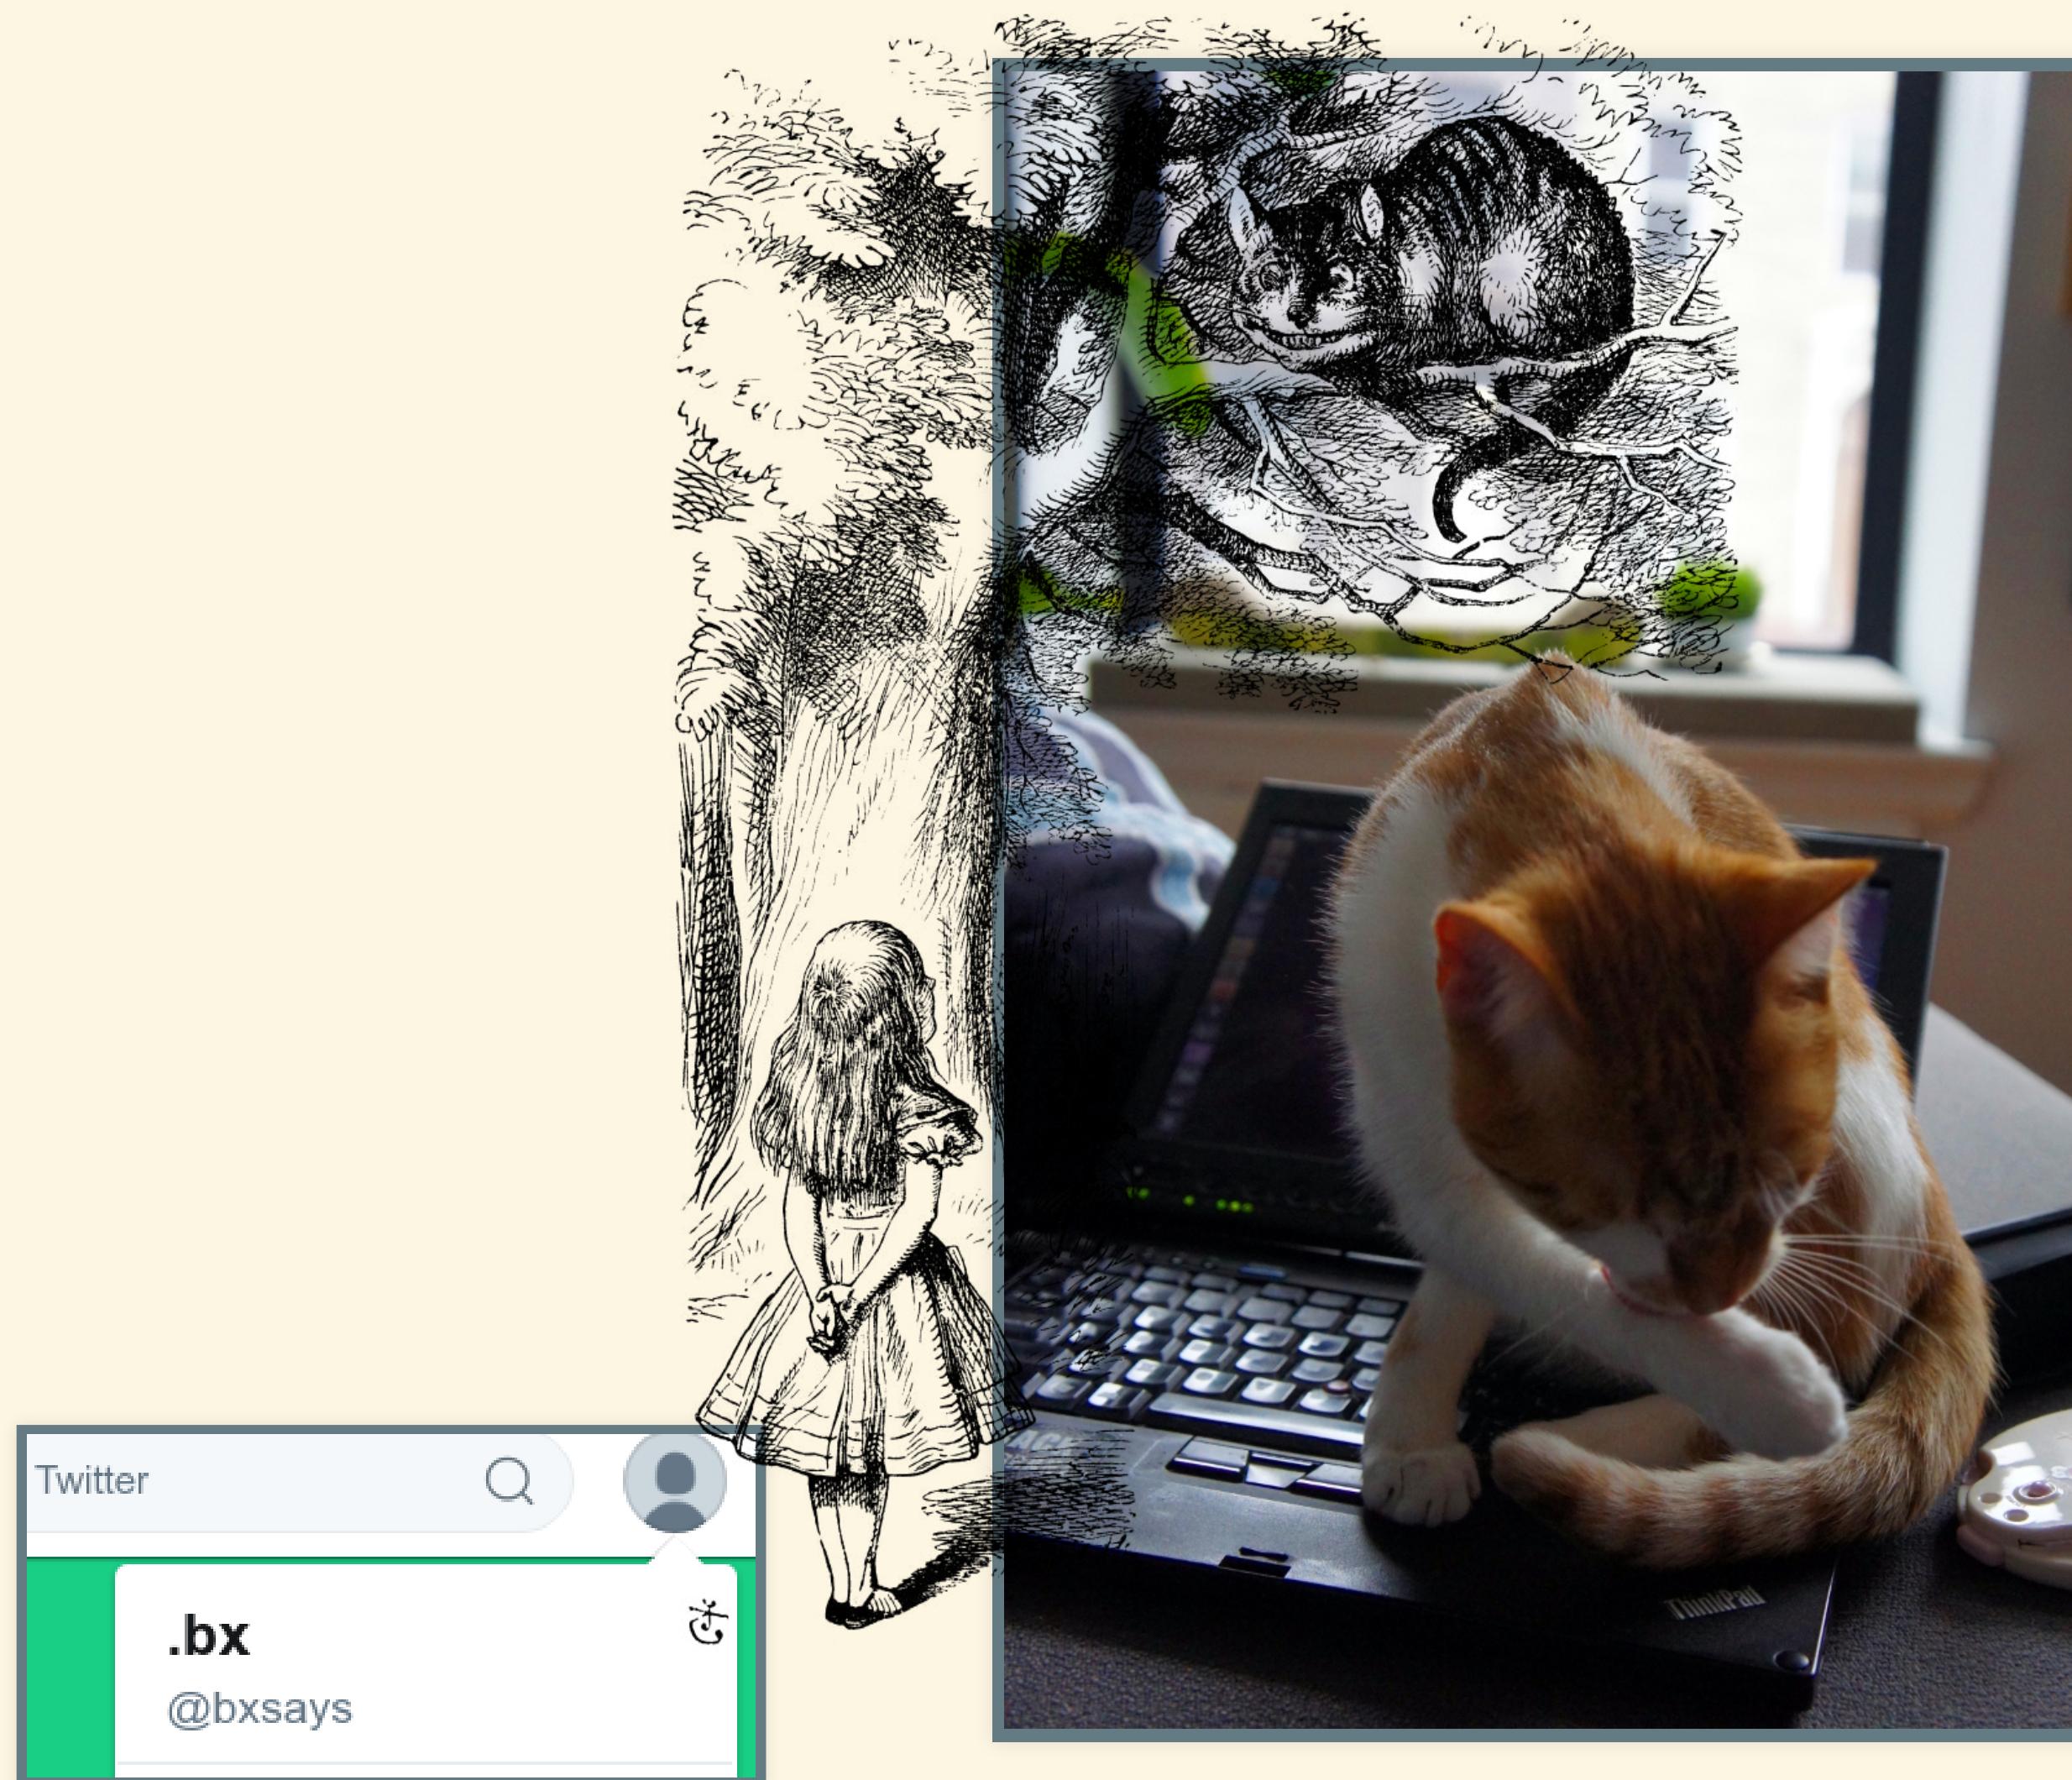


Thank you



More details and tools at: <https://typedregions.com>
Blog post on bootloaders at: <https://bit.ly/2zbentY>
Many thanks **Sergey Bratus**, my PhD advisor

Thank you



bx@narfindustries.com

More details and tools at: <https://typedregions.com>
Blog post on bootloaders at: <https://bit.ly/2zbentY>
Many thanks **Sergey Bratus**, my PhD advisor

References

See also typedregions.com

- [1]. R. .. Shapiro, "The care and feeding of weird machines found in executable metadata," in *Chaos communication congress*, 2012.
- [2]. R. Shapiro, S. Bratus, and S. W. Smith, "Weird machines in ELF: A spotlight on the underappreciated metadata," in *Workshop on offensive technologies*, Washington, D.C., 2013.
- [3]. R. .. Shapiro, "Returning from ELF to libc," in *POC or GTFO*, vol. 0x0, 2013.
- [4]. R. .. Shapiro, "Calling putchar() from an ELF weird machine," in *POC or GTFO*, vol. 0x02, 2013.
- [5]. J. Bangert, S. Bratus, R. Shapiro, and S. W. Smith, "The page-fault weird machine: Lessons in instruction-less computation," in *Workshop on offensive technologies*, Washington, D.C., 2013.
- [6]. J. Bangert, S. Bratus, R. Shapiro, M. E. Locasto, J. Reeves, S. W. Smith, and A. Shubina, "ELFbac: Using the loader format for intent-level semantics and fine-grained protection," Dartmouth College, Computer Science, Hanover, NH, TR2013-727, May 2013.
- [7]. S. Bratus, M. E. Locasto, and M. L. Patterson, "Exploit programming: From buffer overflows to "weird machines" and theory of computation," in *USENIX; login*, 2011, pp. 13-21.
- [8]. S. Bratus and J. Bangert, "ELFs are dorky, elves are cool," in *POC or GTFO*, vol. 0x0, 2013.
- [9]. J. Oakley and S. Bratus, "Exploiting the hard-working DWARF: Trojan and exploit techniques with no native executable code," in *Workshop on offensive technologies*, 2011, p. 11.